

Arithmetic-Logic Unit (ALU)

Arithmetic-Logic Unit: ALU

- *ALU*: Component that can perform various arithmetic (add, subtract, increment, etc.) and logic (AND, OR, etc.) operations, based on control inputs

TABLE 4.2 Desired calculator operations

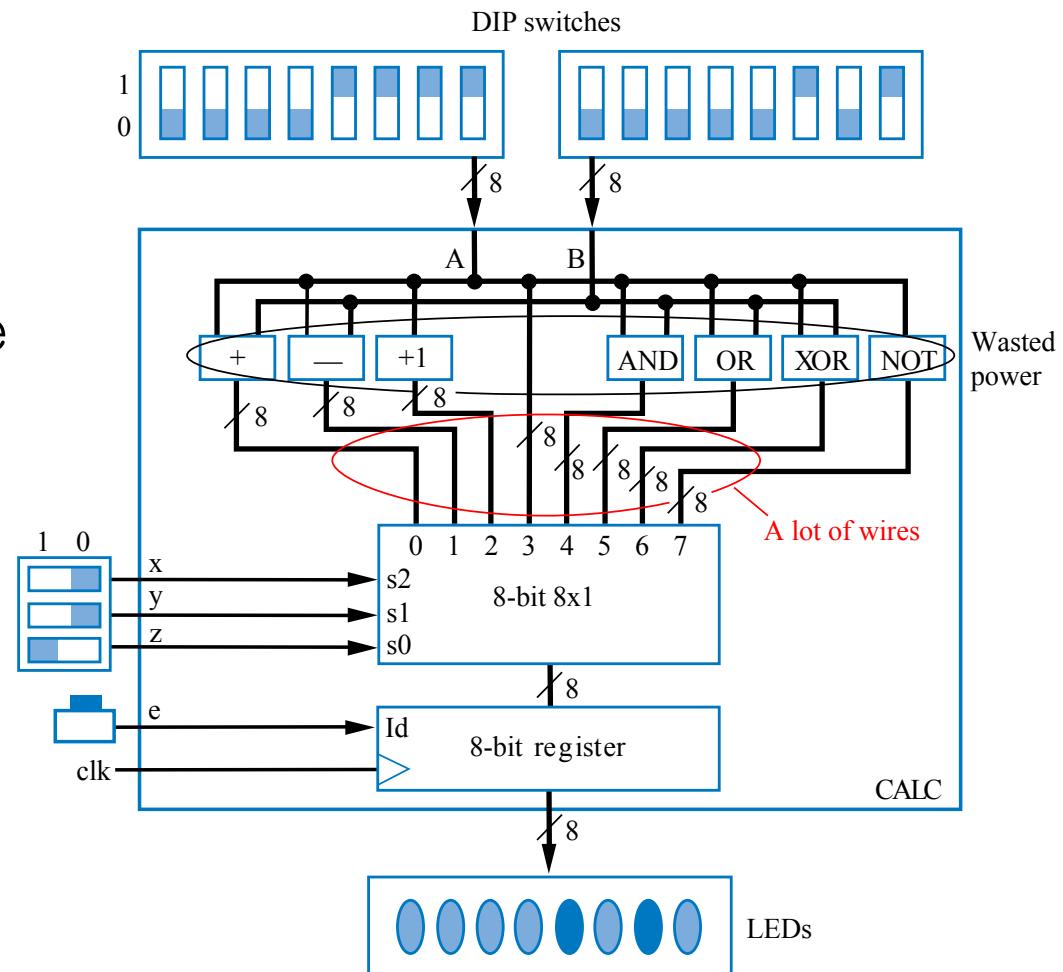
Inputs			Operation	Sample output if A=00001111, B=00000101
x	y	z		
0	0	0	$S = A + B$	S=00010100
0	0	1	$S = A - B$	S=00001010
0	1	0	$S = A + 1$	S=00010000
0	1	1	$S = A$	S=00001111
1	0	0	$S = A \text{ AND } B$ (bitwise AND)	S=00000101
1	0	1	$S = A \text{ OR } B$ (bitwise OR)	S=00001111
1	1	0	$S = A \text{ XOR } B$ (bitwise XOR)	S=00001010
1	1	1	$S = \text{NOT } A$ (bitwise complement)	S=11110000

Multifunction Calculator without an ALU

- Can build using separate components for each operation, and muxes
 - Too many wires, also wastes power computing operations when only use one result at given time

TABLE 4.2 Desired calculator operations

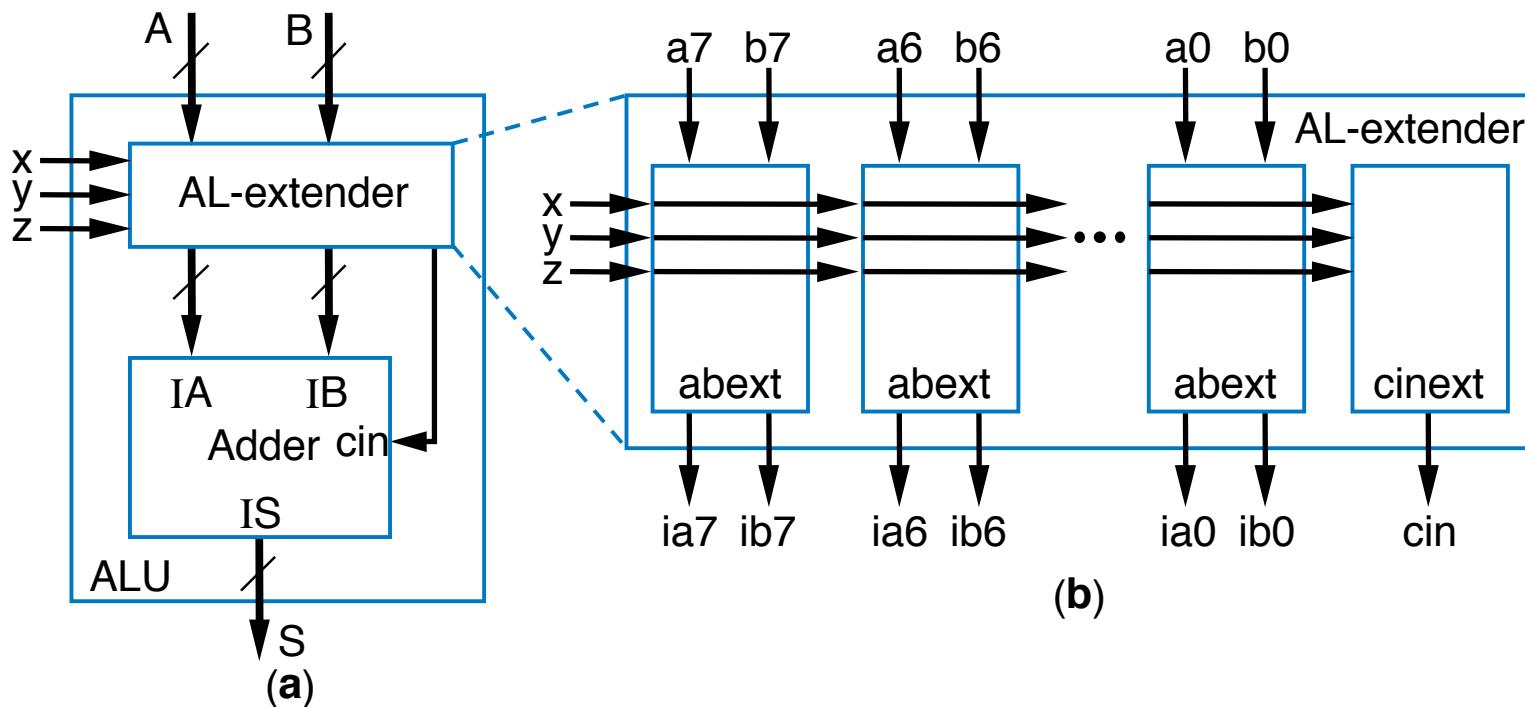
Inputs	Operation	Sample output if A=00001111, B=00000101
x y z		
0 0 0	$S = A + B$	S=00010100
0 0 1	$S = A - B$	S=00001010
0 1 0	$S = A + 1$	S=00010000
0 1 1	$S = A$	S=00001111
1 0 0	$S = A \text{ AND } B$ (bitwise AND)	S=00000101
1 0 1	$S = A \text{ OR } B$ (bitwise OR)	S=00001111
1 1 0	$S = A \text{ XOR } B$ (bitwise XOR)	S=00001010
1 1 1	$S = \text{NOT } A$ (bitwise complement)	S=11110000



ALU

More efficient design uses ALU

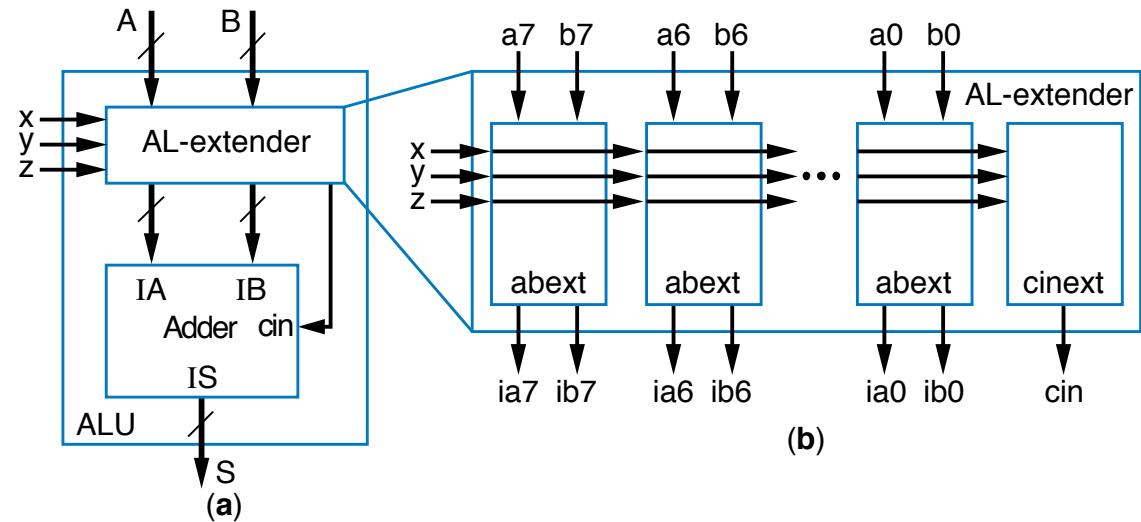
- ALU design not just separate components multiplexed (same problem as previous slide)
- Instead, ALU design uses single adder, plus logic in front of adder's A and B inputs
 - Logic in front is called an *arithmetic-logic extender*
- Extender modifies A and B inputs so desired operation appears at output of the adder



Arithmetic-Logic Extender in Front of ALU

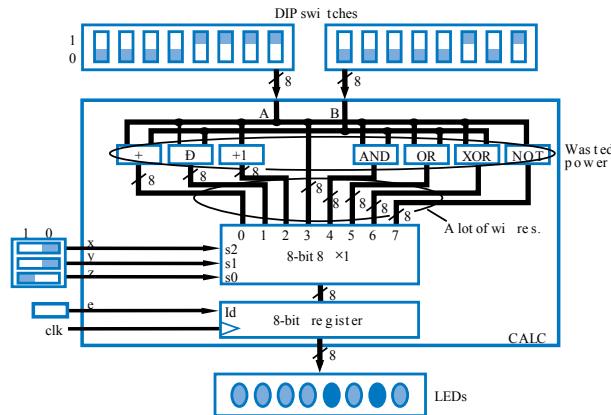
TABLE 4.2 Desired calculator operations

Inputs			Operation	Sample output if A=00001111, B=00000101
x	y	z		
0	0	0	S = A + B	S=00010100
0	0	1	S = A - B	S=00001010
0	1	0	S = A + 1	S=00010000
0	1	1	S = A	S=00001111
1	0	0	S = A AND B (bitwise AND)	S=00000101
1	0	1	S = A OR B (bitwise OR)	S=00001111
1	1	0	S = A XOR B (bitwise XOR)	S=00001010
1	1	1	S = NOT A (bitwise complement)	S=11110000

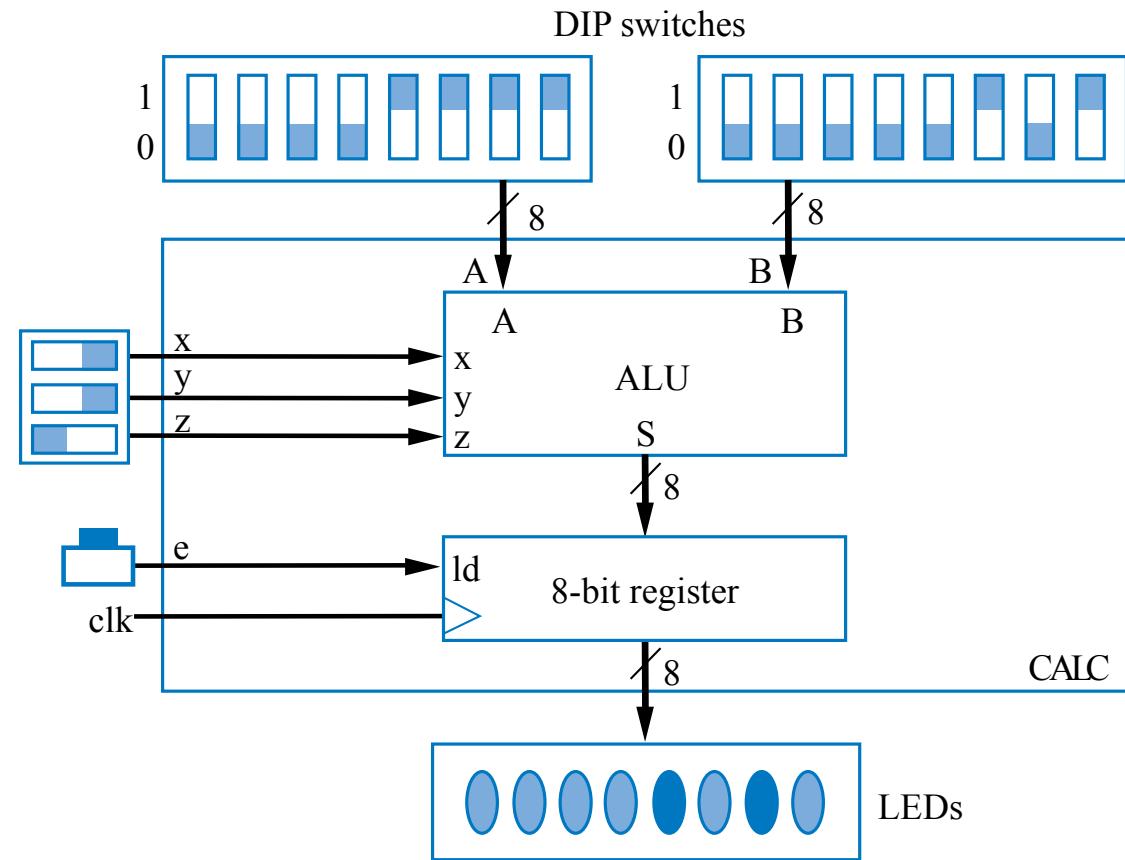


- xyz=000 Want S=A+B : just pass a to ia, b to ib, and set cin=0
- xyz=001 Want S=A-B : pass a to ia, b' to ib and set cin=1 (two's complement)
- xyz=010 Want S=A+1 : pass a to ia, set ib=0, and set cin=1
- xyz=011 Want S=A : pass a to ia, set ib=0, and set cin=0
- xyz=100 Want S=A AND B : set ia=a*b, b=0, and cin=0
- Others: likewise
- Based on above, create logic for ia(x,y,z,a,b) and ib(x,y,z,a,b) for each abext, and create logic for cin(x,y,z), to complete design of the AL-extender component

ALU Example: Multifunction Calculator



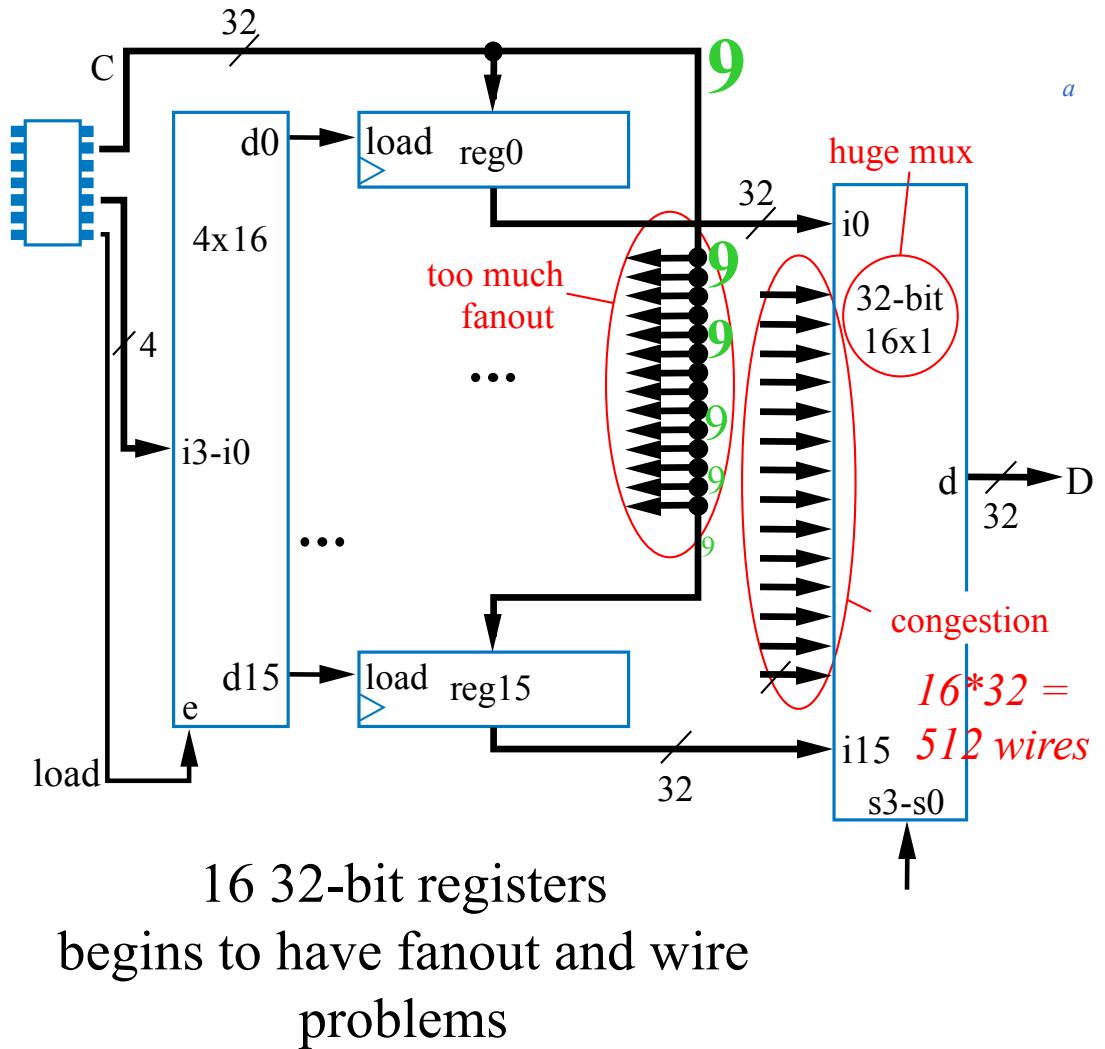
- Design using ALU is elegant and efficient
 - No mass of wires
 - No big waste of power



Register Files (Register Bank)

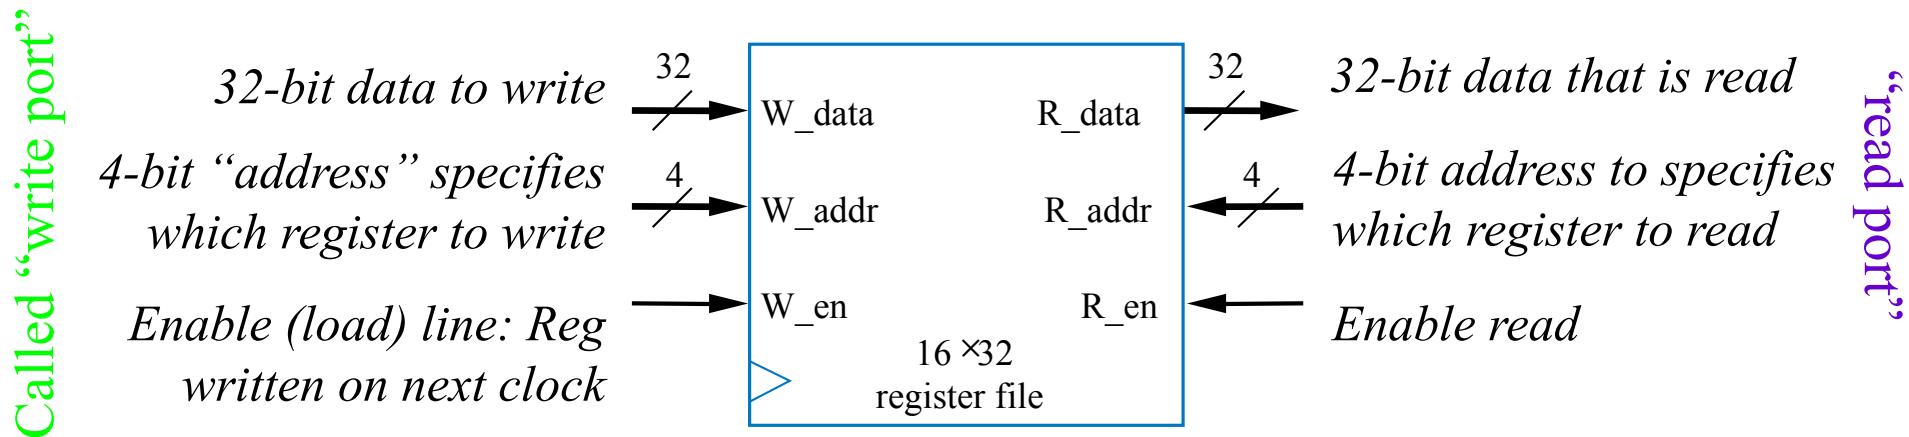
Register Files

- Accessing one of several registers is:
 - OK if just a few registers
 - Problematic when many
 - Ex: Earlier above-mirror display, with 16 registers
 - Much fanout (branching of wire): Weakens signal
 - Many wires: Congestion



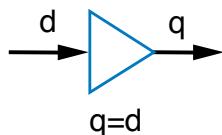
Register File

- **MxN register file:** Efficient design for one-at-a-time write/read of many registers
 - Consider 16 32-bit registers



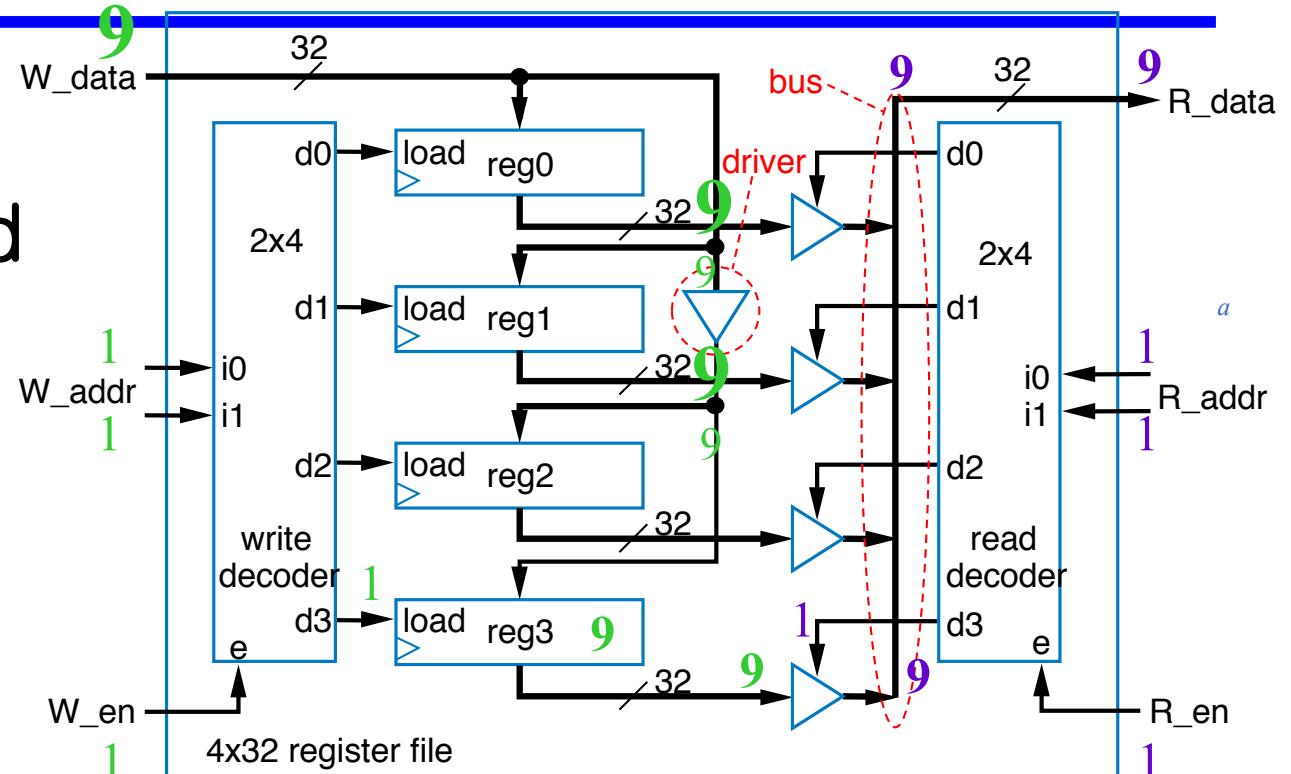
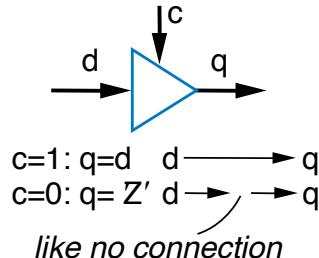
Register File

- Internal design uses drivers and bus driver

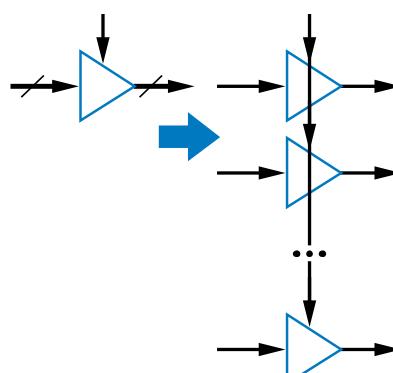


Boosts signal

three-state driver



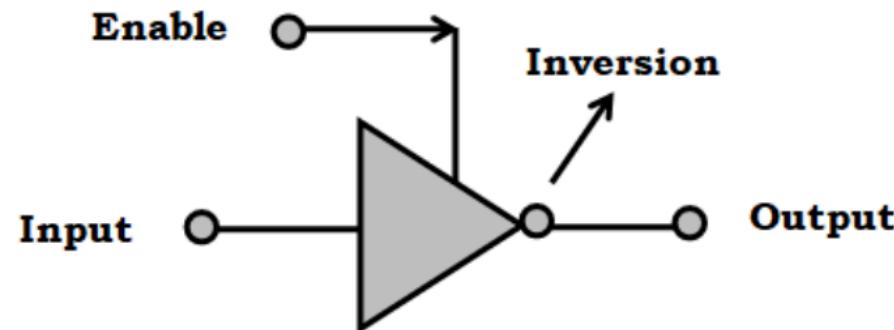
Internal design of 4x32 RF; 16x32 RF follows similarly



Note: Each driver in figure actually represents 32 1-bit drivers

Tri-state buffer

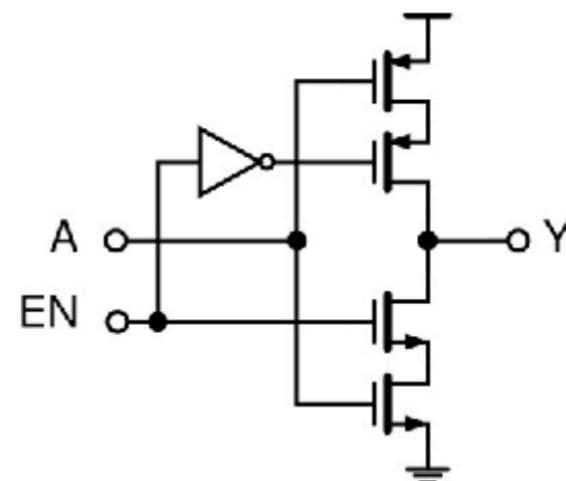
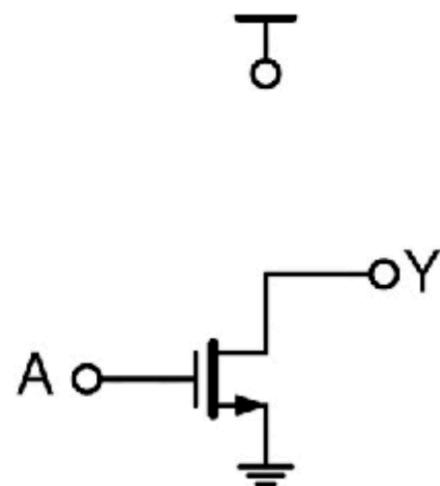
- Connect the input to output when enabled
- When disabled,
 - Output is not connected to any source
 - It is unstable state, and different from logical '0'



Enable	Input	Output
0	0	Hi-Z
0	1	Hi-Z
1	0	1
1	1	0

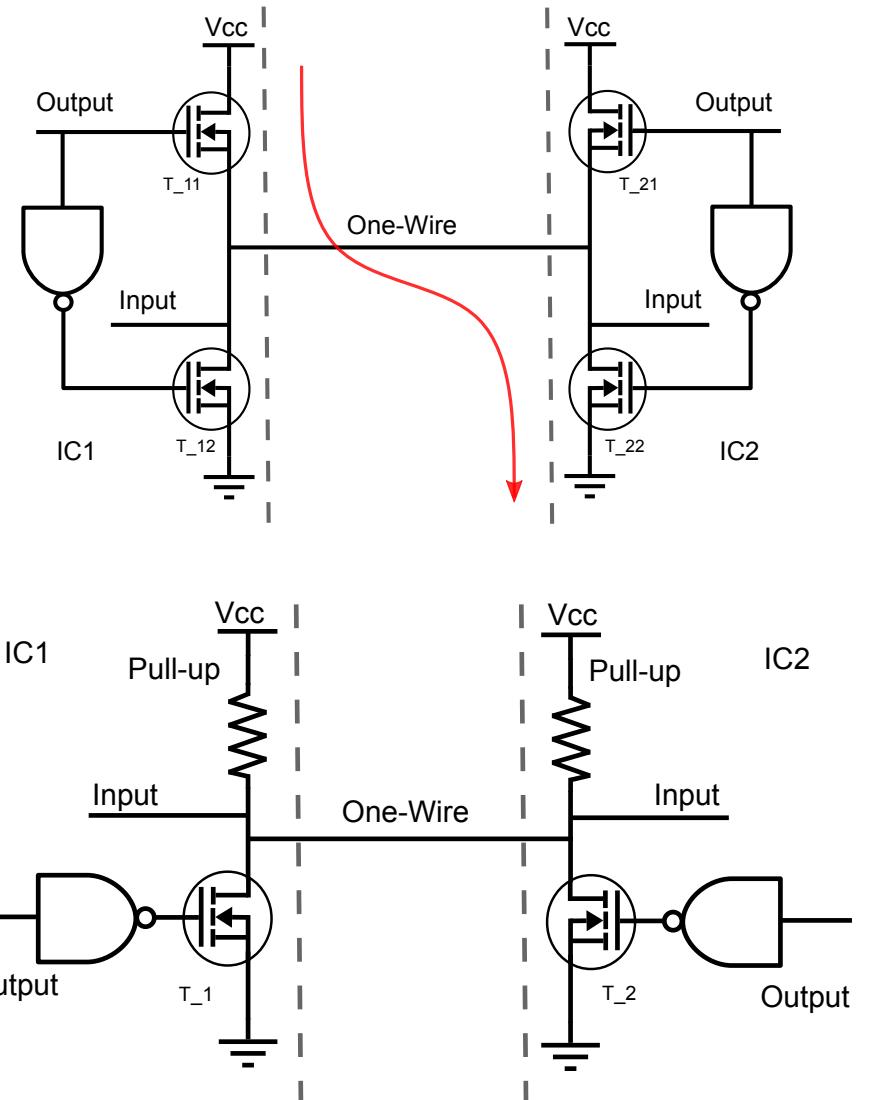
Tri-state buffer

- High-Z (High impedance)
 - High enough impedance (resistance) to be regarded as disconnected



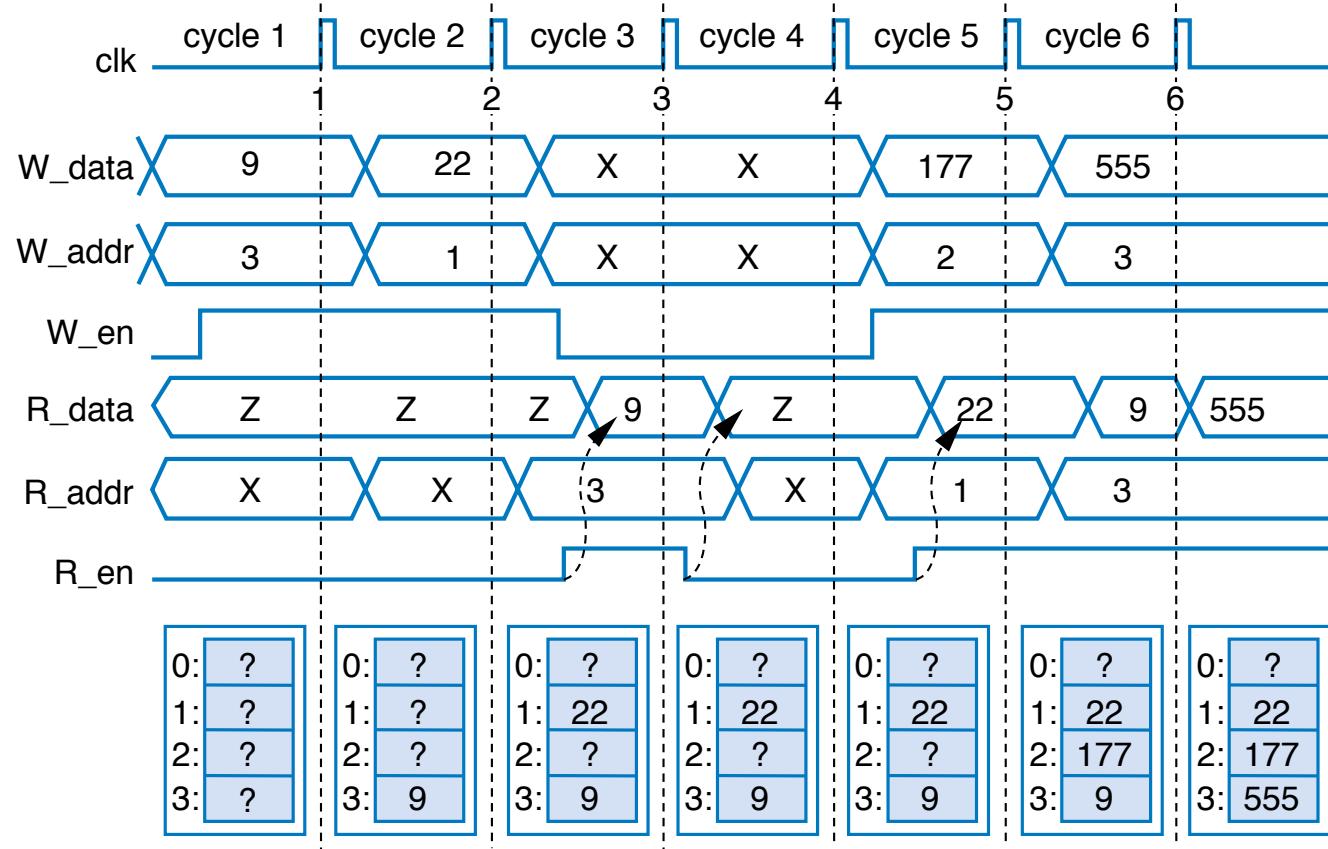
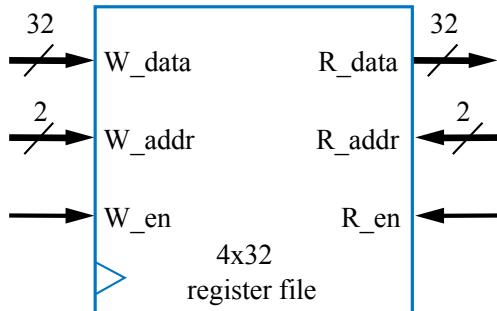
Open drain driver

- Bus multiplexer
 - Aggregate output from multiple sources
- When multiple CMOS drivers are connected to one node,
 - Logical '0' and '1' values may exist at the same time
 - It causes logical conflict, and electrical short circuit
- How to avoid it?
 - Remove one side of driver capability.



Register File Timing Diagram

- Can write one register and read one register each clock cycle
 - May be same register

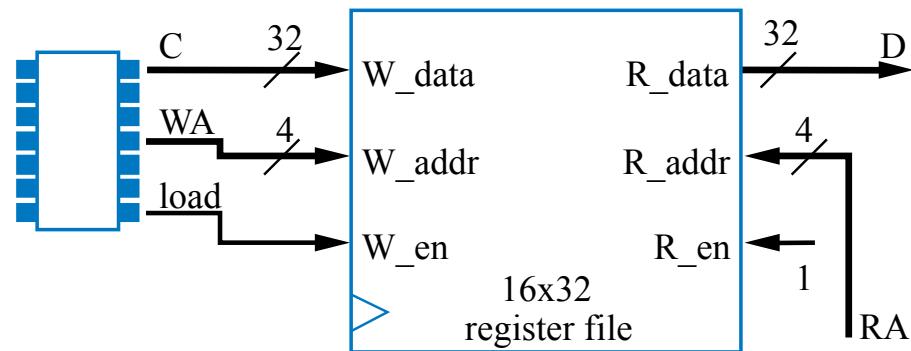


Chapter Summary

- Need datapath components to store and operate on multi-bit data
 - Also known as register-transfer-level (RTL) components
- Components introduced
 - Registers
 - Adders
 - Comparators
 - Multipliers
 - Subtractors
 - Arithmetic–Logic Units
 - Shifters
 - Counters and Timers
 - Register Files
- Next chapter combines knowledge of combinational logic design, sequential logic design, and datapath components, to build digital circuits that can perform general and powerful computations

Register-File Example: Above-Mirror Display

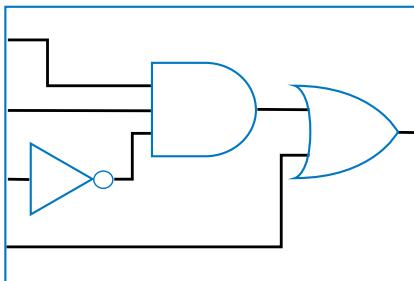
- 16 32-bit registers that can be written by car's computer, and displayed
 - Use 16x32 register file
 - Simple, elegant design
- Register file hides complexity internally
 - And because only one register needs to be written and/or read at a time, internal design is simple



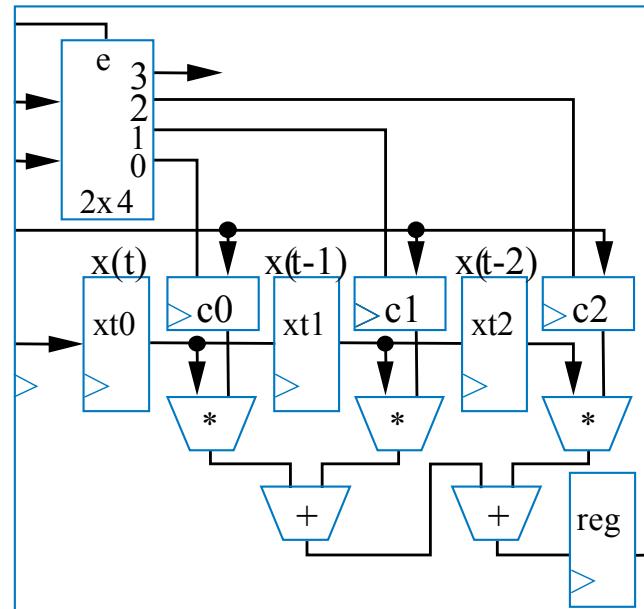
Programmable processor

Introduction

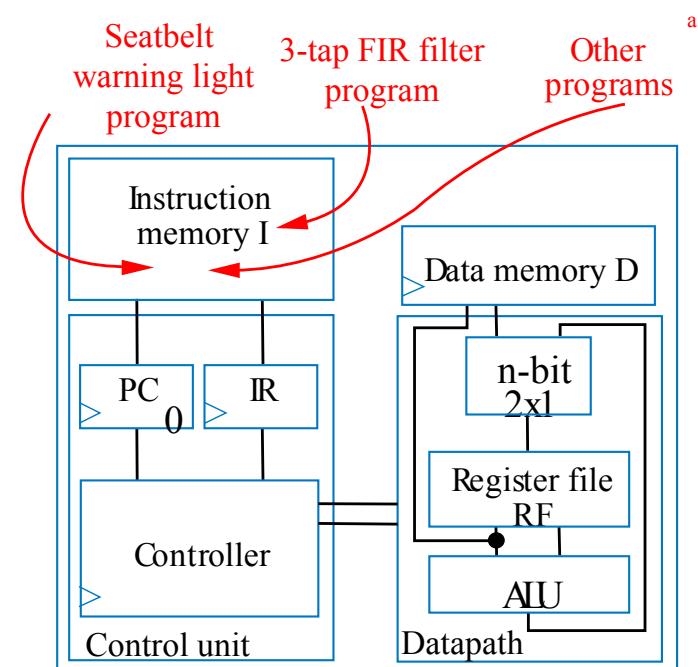
- Programmable (general-purpose) processor
 - Mass-produced, then programmed to implement different processing tasks
 - Well-known common programmable processors: Pentium, Sparc, PowerPC
 - Lesser-known but still common: ARM, MIPS, 8051, PIC, AVR
 - Low-cost embedded processors found in cell phones, blinking shoes, etc.
 - Instructive to design a very simple programmable processor
 - Real processors can be much more complex



Seatbelt warning light single-purpose processor



3-tap FIR filter single-purpose processor



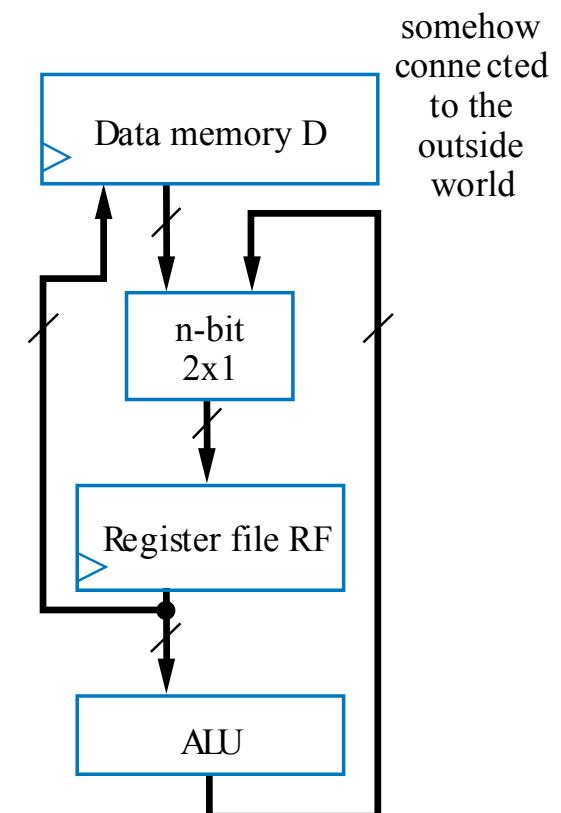
General-purpose processor

2



Basic Architecture

- Processing generally consists of:
 - Loading some data
 - Transforming that data
 - Storing that data
- *Basic datapath*: Useful circuit in a programmable processor
 - Can read/write data memory, where main data exists
 - Has register file to hold data locally
 - Has ALU to transform local data

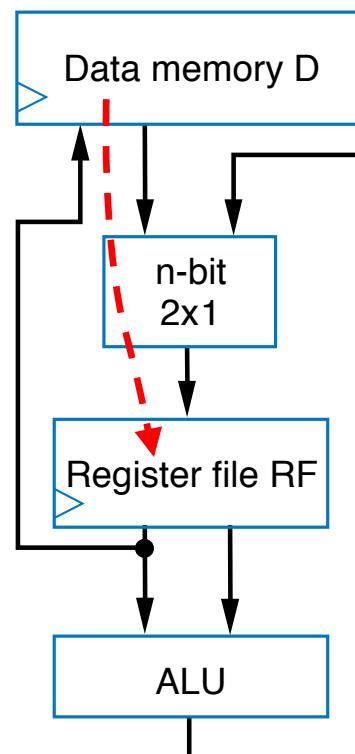


Datapath

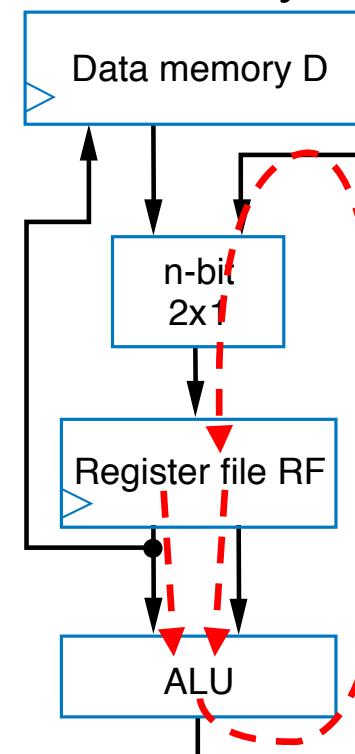


Basic Datapath Operations

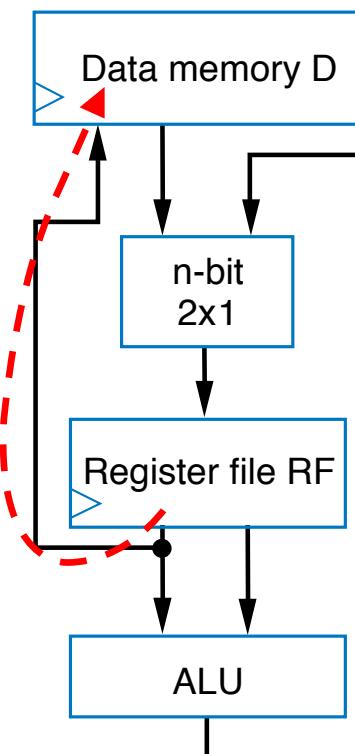
- Load operation: Load data from data memory to RF
- ALU operation: Transforms data by passing one or two RF register values through ALU, performing operation (ADD, SUB, AND, OR, etc.), and writing back into RF.
- Store operation: Stores RF register value back into data memory
- Each operation can be done in one clock cycle



Load operation



ALU operation

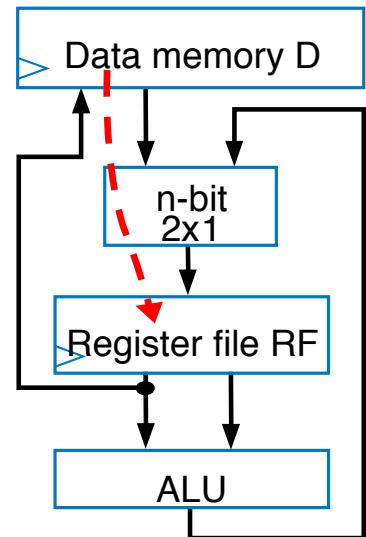


Store operation

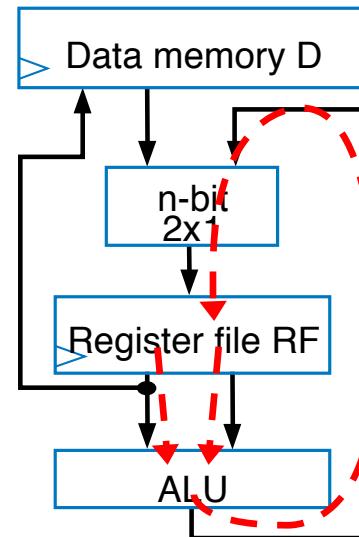


Basic Datapath Operations

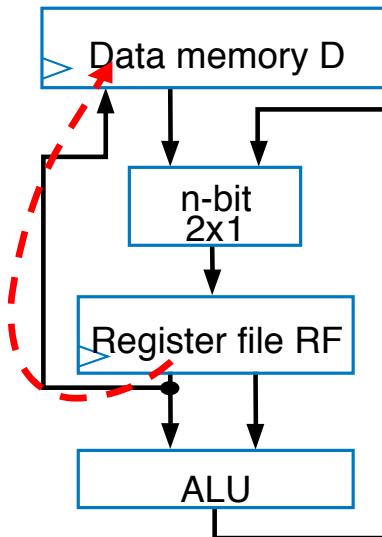
- **Q:** Which are valid *single-cycle operations* for given datapath?
 - Move D[1] to RF[1] (i.e., RF[1] = D[1])
 - **A:** YES – That's a load operation
 - Store RF[1] to D[9] and store RF[2] to D[10]
 - **A:** NO – Requires two separate store operations
 - Add D[0] plus D[1], store result in D[9]
 - **A:** NO – ALU operation (ADD) only works with RF. Requires two load operations (e.g., RF[0]=D[0]; RF[1]=D[1], an ALU operation (e.g., RF[2]=RF[0]+RF[1]), and a store operation (e.g., D[9]=RF[2])



Load operation



ALU operation

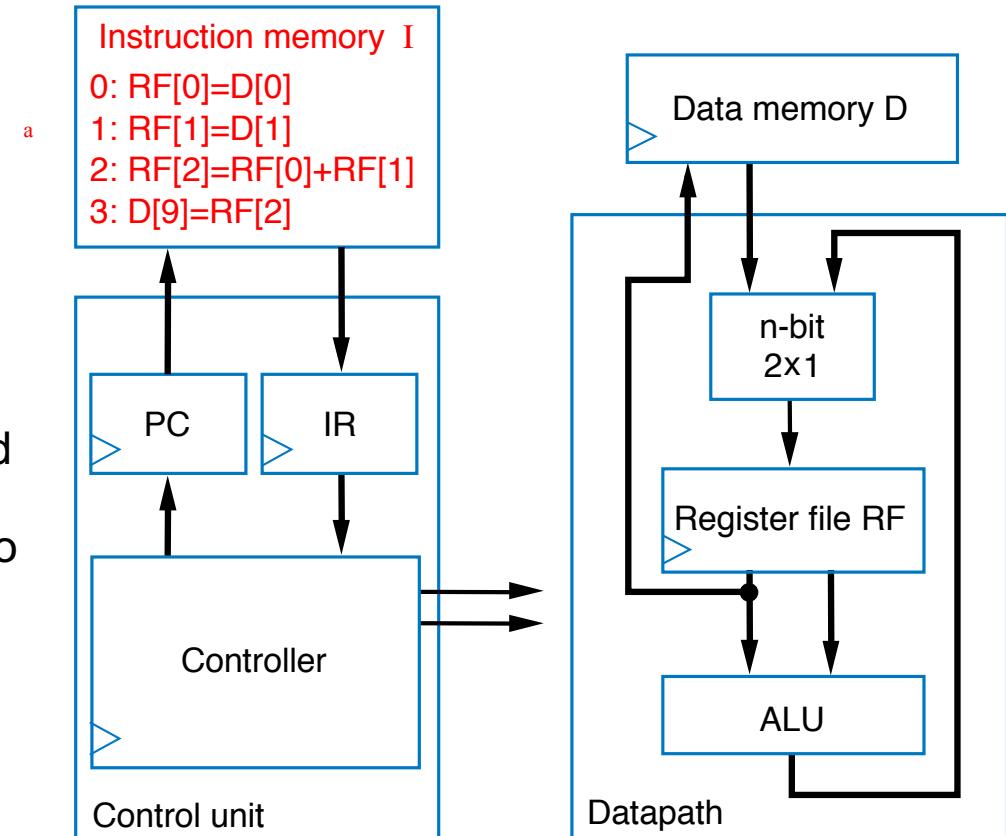


Store operation



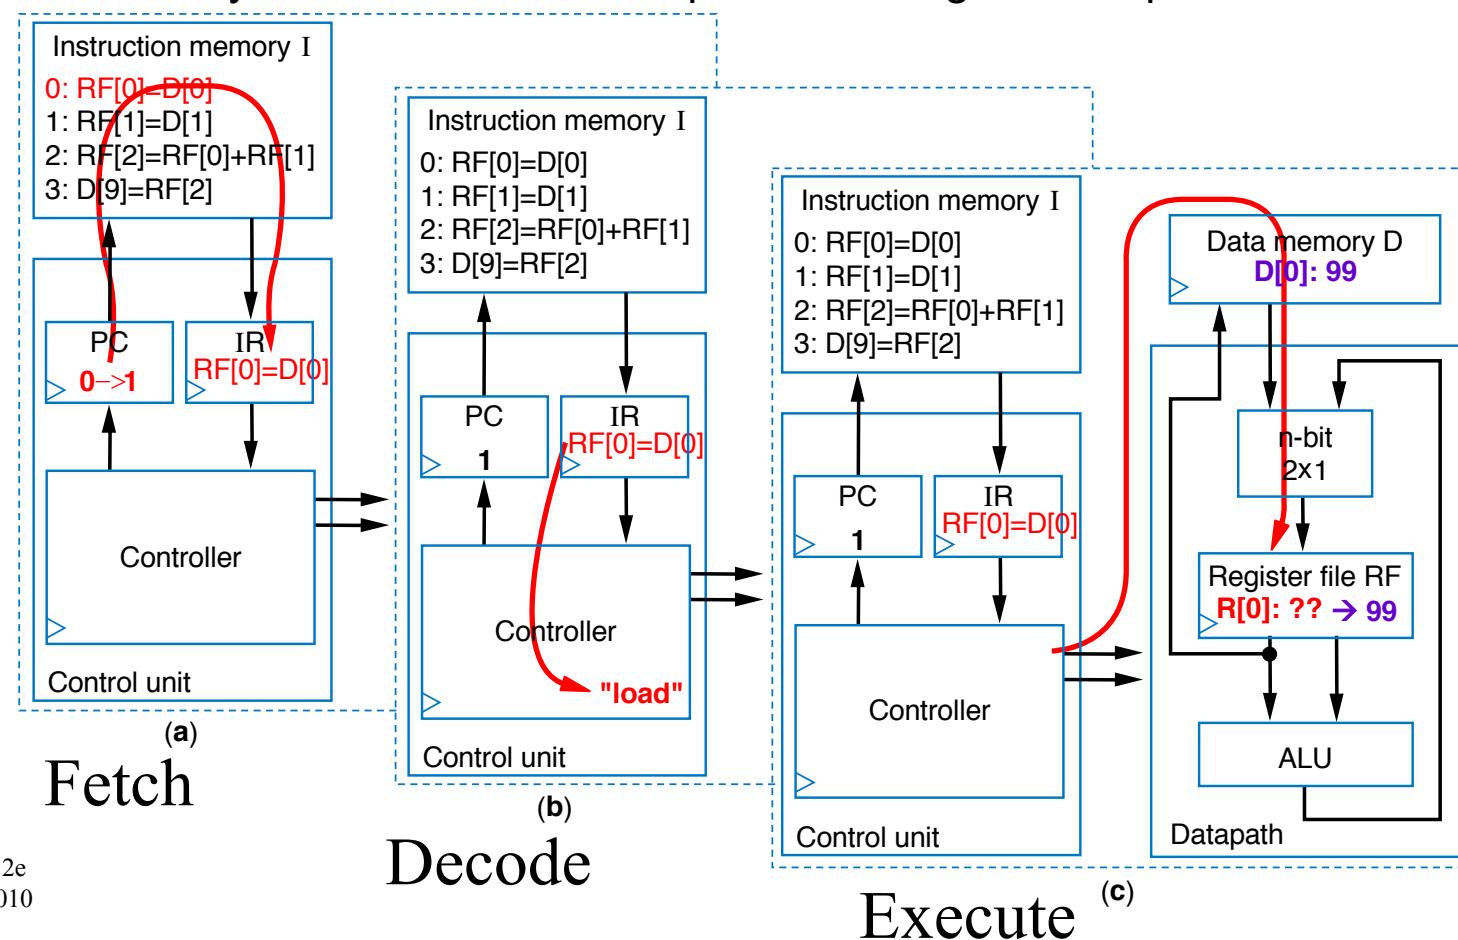
Basic Architecture – Control Unit

- $D[9] = D[0] + D[1]$ – requires a sequence of four datapath operations:
 - 0: $RF[0] = D[0]$
 - 1: $RF[1] = D[1]$
 - 2: $RF[2] = RF[0] + RF[1]$
 - 3: $D[9] = RF[2]$
- Each operation is an *instruction*
 - Sequence of instructions – *program*
 - Looks cumbersome, but that's the world of programmable processors – Decomposing desired computations into processor-supported operations
 - Store program in *Instruction memory*
 - *Control unit* reads each instruction and executes it on the datapath
 - PC: Program counter – address of current instruction
 - IR: Instruction register – current instruction



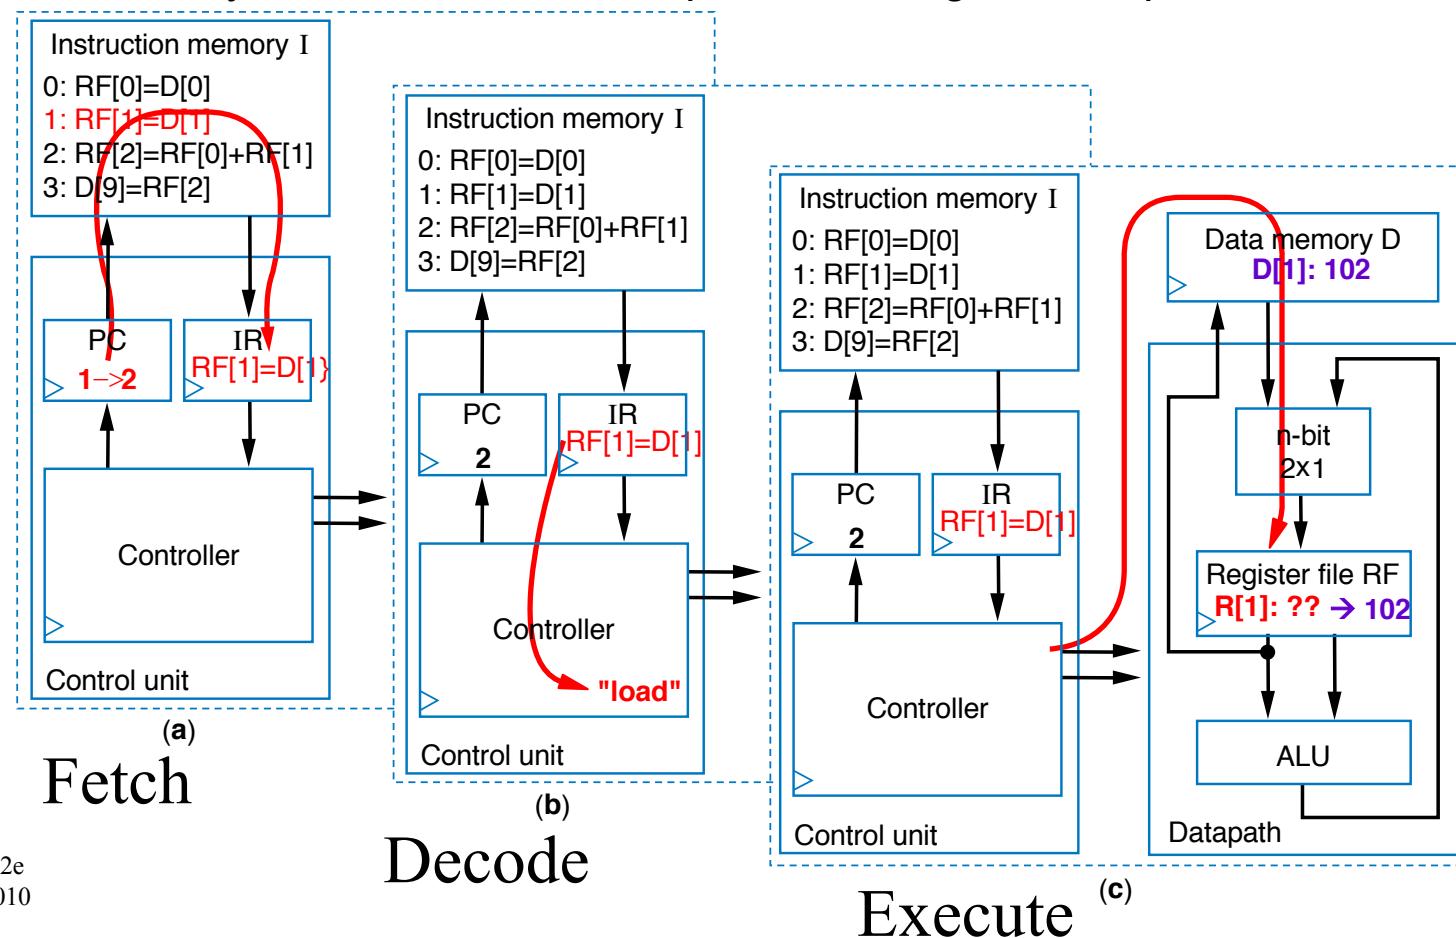
Basic Architecture – Control Unit

- To carry out each *instruction*, the control unit must:
 - Fetch – Read instruction from inst. mem.
 - Decode – Determine the operation and operands of the instruction
 - Execute – Carry out the instruction's operation using the datapath



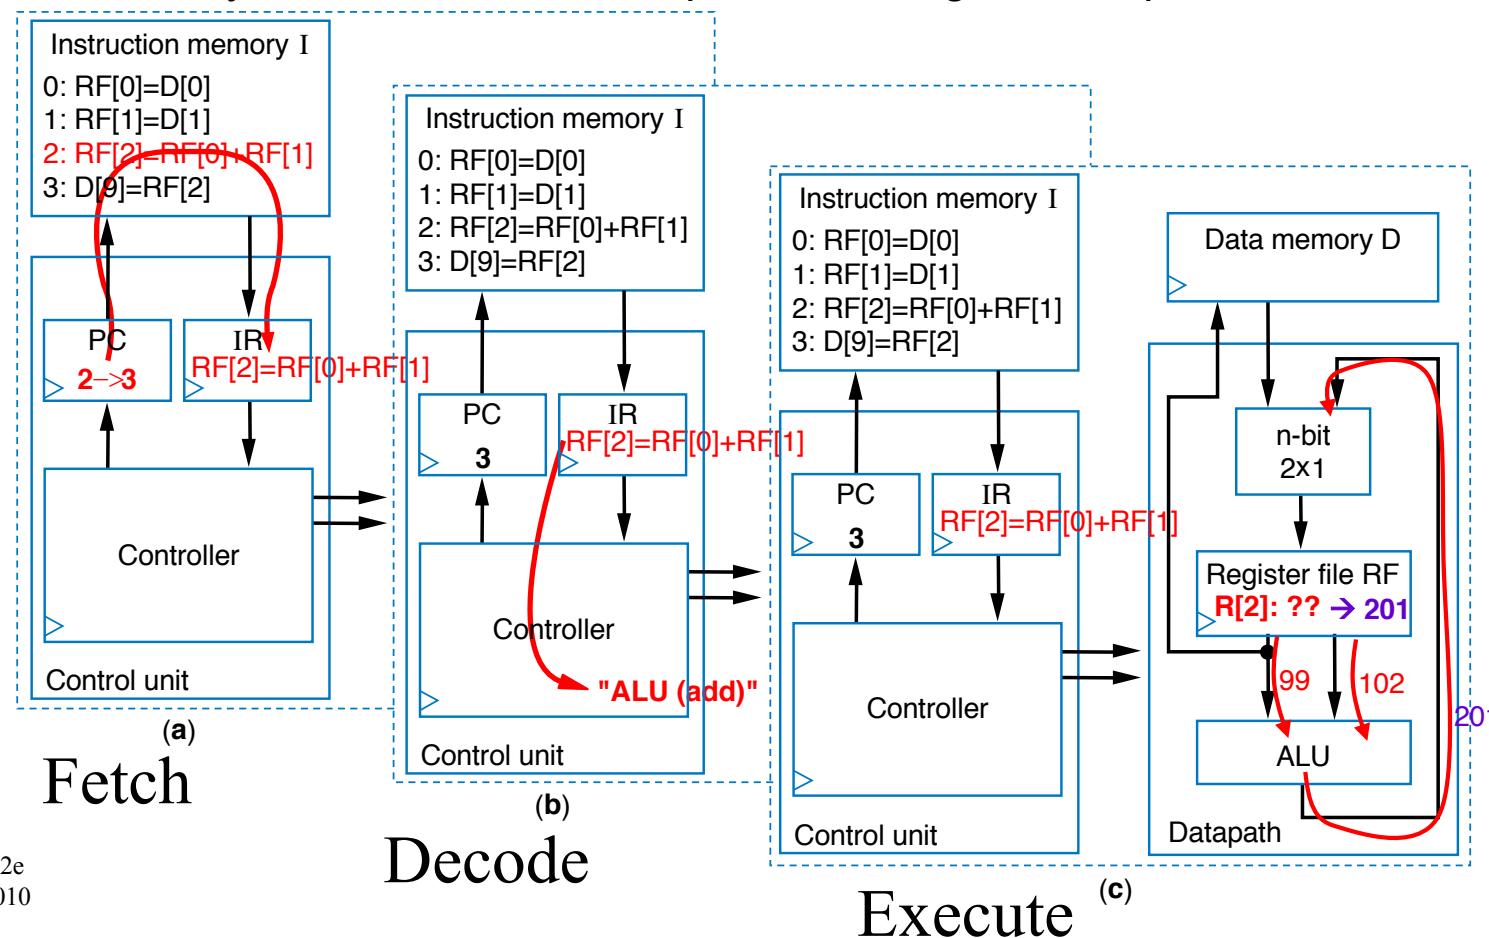
Basic Architecture – Control Unit

- To carry out each *instruction*, the control unit must:
 - Fetch – Read instruction from inst. mem.
 - Decode – Determine the operation and operands of the instruction
 - Execute – Carry out the instruction's operation using the datapath



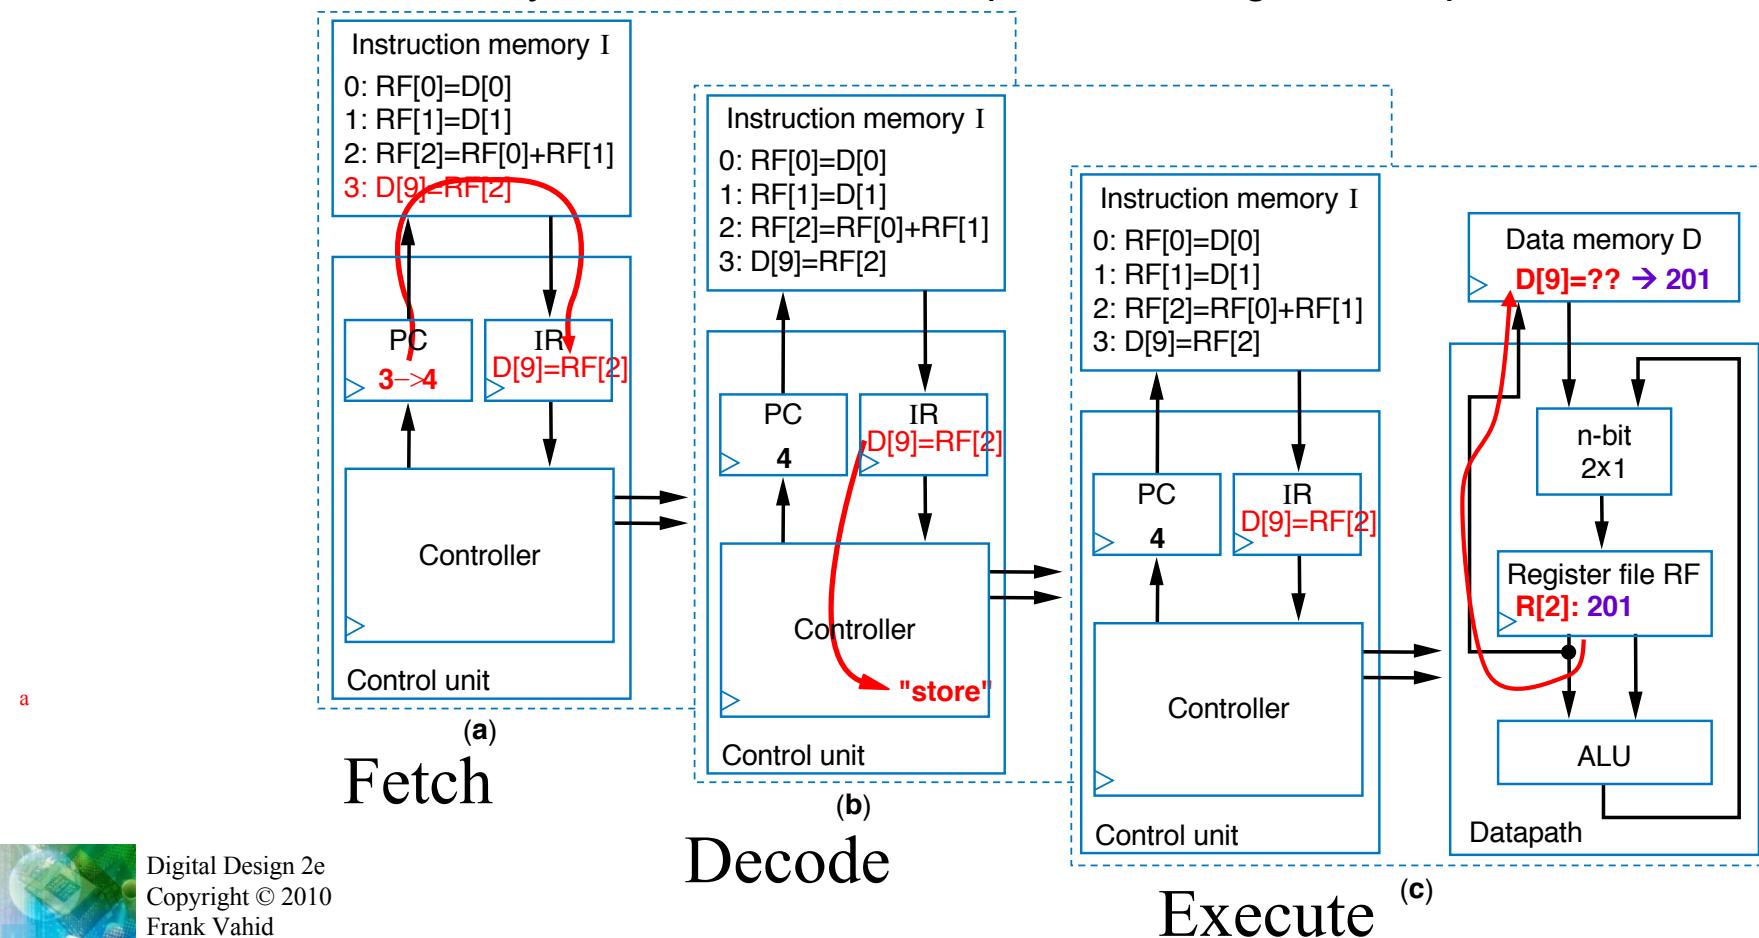
Basic Architecture – Control Unit

- To carry out each *instruction*, the control unit must:
 - Fetch – Read instruction from inst. mem.
 - Decode – Determine the operation and operands of the instruction
 - Execute – Carry out the instruction's operation using the datapath



Basic Architecture – Control Unit

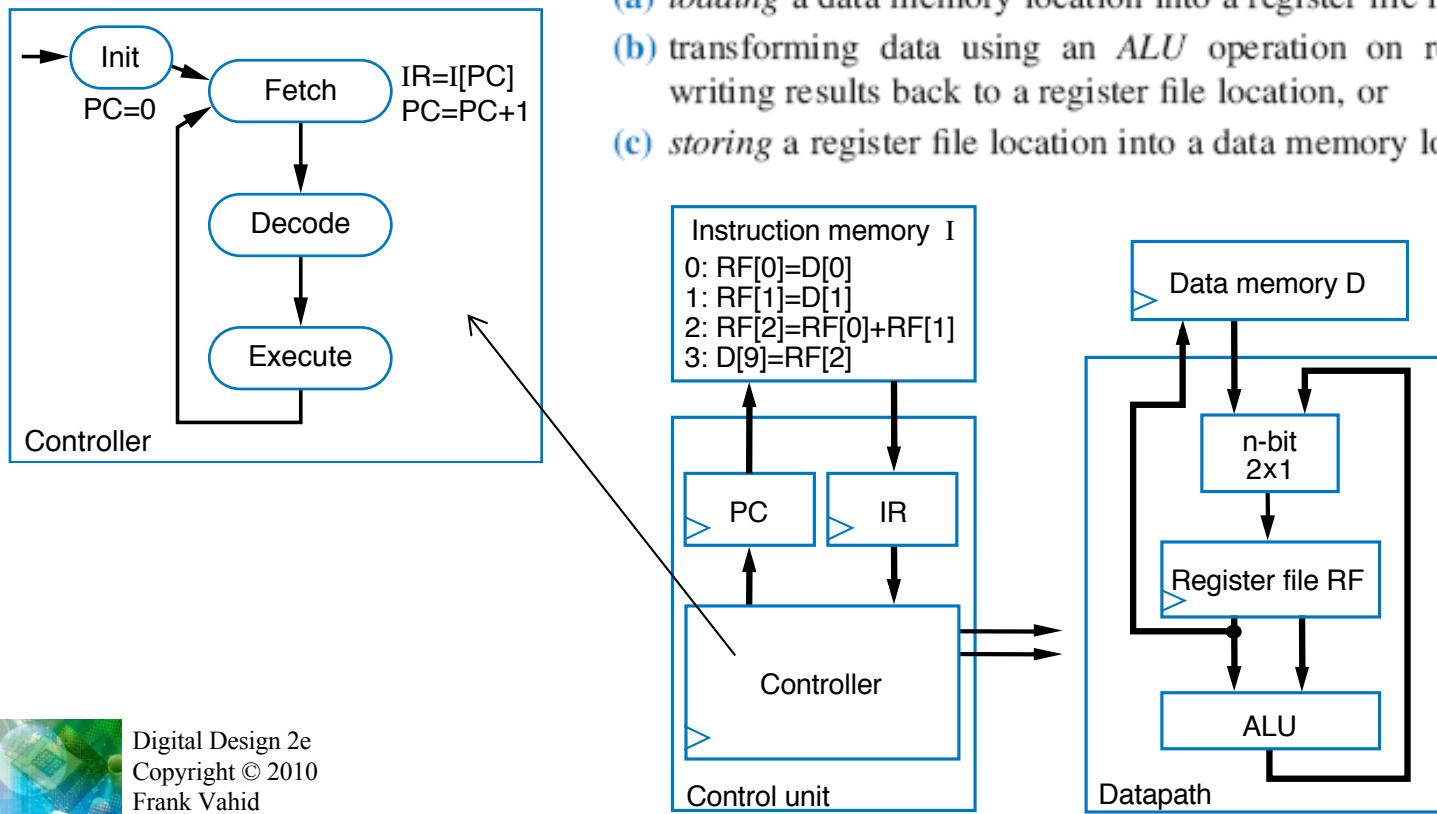
- To carry out each *instruction*, the control unit must:
 - Fetch – Read instruction from inst. mem.
 - Decode – Determine the operation and operands of the instruction
 - Execute – Carry out the instruction's operation using the datapath



Basic Architecture – Control Unit

To summarize, the control unit processes each instruction in three stages:

1. first *fetching* the instruction by loading the current instruction into *IR* and incrementing the *PC* for the next fetch,
2. next *decoding* the instruction to determine its operation, and
3. finally *executing* the operation by setting the appropriate control lines for the datapath, if applicable. If the operation is a datapath operation, the operation may be one of three possible types:
 - (a) *loading* a data memory location into a register file location,
 - (b) transforming data using an *ALU* operation on register file locations and writing results back to a register file location, or
 - (c) *storing* a register file location into a data memory location.



Creating a Sequence of Instructions

- **Q:** Create sequence of instructions to compute $D[3] = D[0]+D[1]+D[2]$ on earlier-introduced processor
- **A1:** One possible sequence
 - First load data memory locations into register file
 - $R[3] = D[0]$
 - $R[4] = D[1]$
 - $R[2] = D[2]$
 - *(Note arbitrary register locations)*
 - Next, perform the additions
 - $R[1] = R[3] + R[4]$
 - $R[1] = R[1] + R[2]$
 - Finally, store result
 - $D[3] = R[1]$
- **A2:** Alternative sequence
 - First load $D[0]$ and $D[1]$ and add them
 - $R[1] = D[0]$
 - $R[2] = D[1]$
 - $R[1] = R[1] + R[2]$
 - Next, load $D[2]$ and add
 - $R[2] = D[2]$
 - $R[1] = R[1] + R[2]$
 - Finally, store result
 - $D[3] = R[1]$



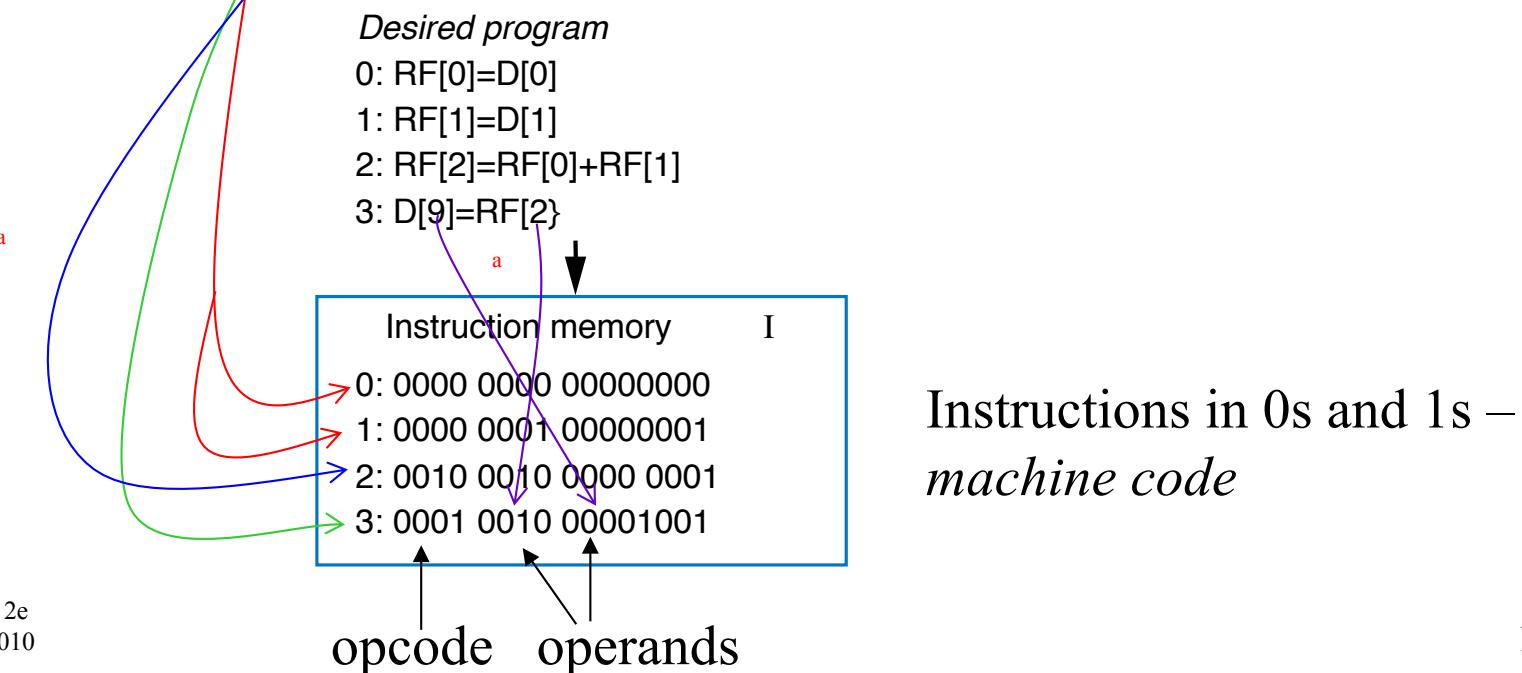
Number of Cycles

- **Q:** How many cycles are needed to execute six instructions using the earlier-described processor?
- **A:** Each instruction requires 3 cycles – 1 to fetch, 1 to decode, and 1 to execute
 - Thus, $6 \text{ instr} * 3 \text{ cycles/instr} = 18 \text{ cycles}$

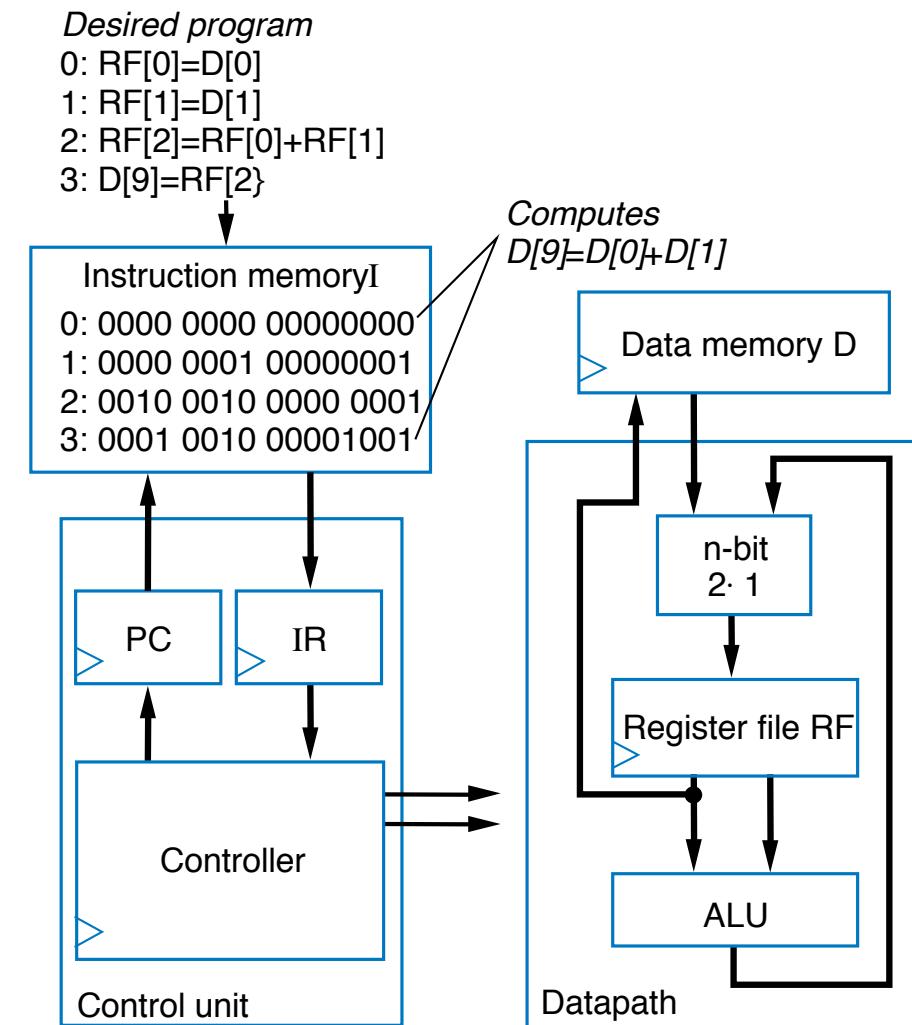


Three-Instruction Programmable Processor

- Instruction Set – List of allowable instructions and their representation in memory, e.g.,
 - Load* instruction – $0000\ r_3r_2r_1r_0\ d_7d_6d_5d_4d_3d_2d_1d_0$
 - Store* instruction – $0001\ r_3r_2r_1r_0\ d_7d_6d_5d_4d_3d_2d_1d_0$
 - Add* instruction – $0010\ ra_3ra_2ra_1ra_0\ rb_3rb_2rb_1rb_0\ rc_3rc_2rc_1rc_0$



Program for Three-Instruction Processor



Program for Three-Instruction Processor

- Another example program in machine code
 - Compute $D[5] = D[5] + D[6] + D[7]$

```
0: 0000 0000 00000101 // RF[0] = D[5]
1: 0000 0001 00000110 // RF[1] = D[6]
2: 0000 0010 00000111 // RF[2] = D[7]
3: 0010 0000 0000 0001 // RF[0] = RF[0] + RF[1]
                         // which is D[5]+D[6]
4: 0010 0000 0000 0010 // RF[0] = RF[0] + RF[2]
                         // now D[5]+D[6]+D[7]
5: 0001 0000 00000101 // D[5] = RF[0]
```

- Load*** instruction—**0000 r₃r₂r₁r₀ d₇d₆d₅d₄d₃d₂d₁d₀**
- Store*** instruction—**0001 r₃r₂r₁r₀ d₇d₆d₅d₄d₃d₂d₁d₀**
- Add*** instruction—**0010 ra₃ra₂ra₁ra₀ rb₃rb₂rb₁rb₀ rc₃rc₂rc₁rc₀**



Assembly Code

- Machine code (0s and 1s) hard to work with
- Assembly code – Uses mnemonics
 - *Load* instruction—**MOV Ra, d**
 - specifies the operation $RF[a]=D[d]$. a must be 0,1,..., or 15—so $R0$ means $RF[0]$, $R1$ means $RF[1]$, etc. d must be 0, 1, ..., 255
 - • *Store* instruction—**MOV d, Ra**
 - specifies the operation $D[d]=RF[a]$
 - • *Add* instruction—**ADD Ra, Rb, Rc**
 - specifies the operation $RF[a]=RF[b]+RF[c]$

Desired program

0: $RF[0]=D[0]$

1: $RF[1]=D[1]$

2: $RF[2]=RF[0]+RF[1]$

3: $D[9]=RF[2]$

0: 0000 0000 00000000

1: 0000 0001 00000001

2: 0010 0010 0000 0001

3: 0001 0010 00001001

0: **MOV R0, 0**

1: **MOV R1, 1**

2: **ADD R2, R0, R1**

3: **MOV 9, R2**

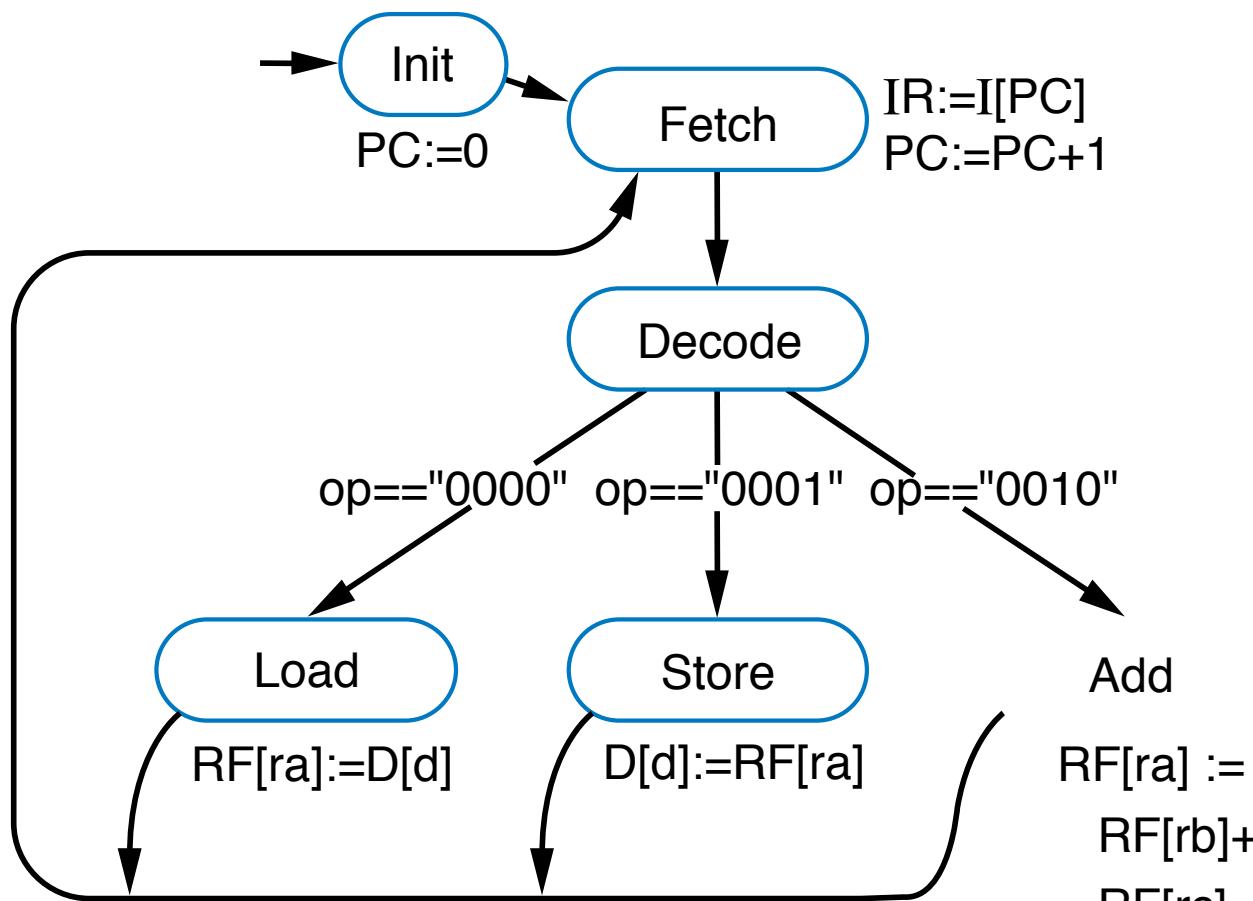
machine code

assembly code



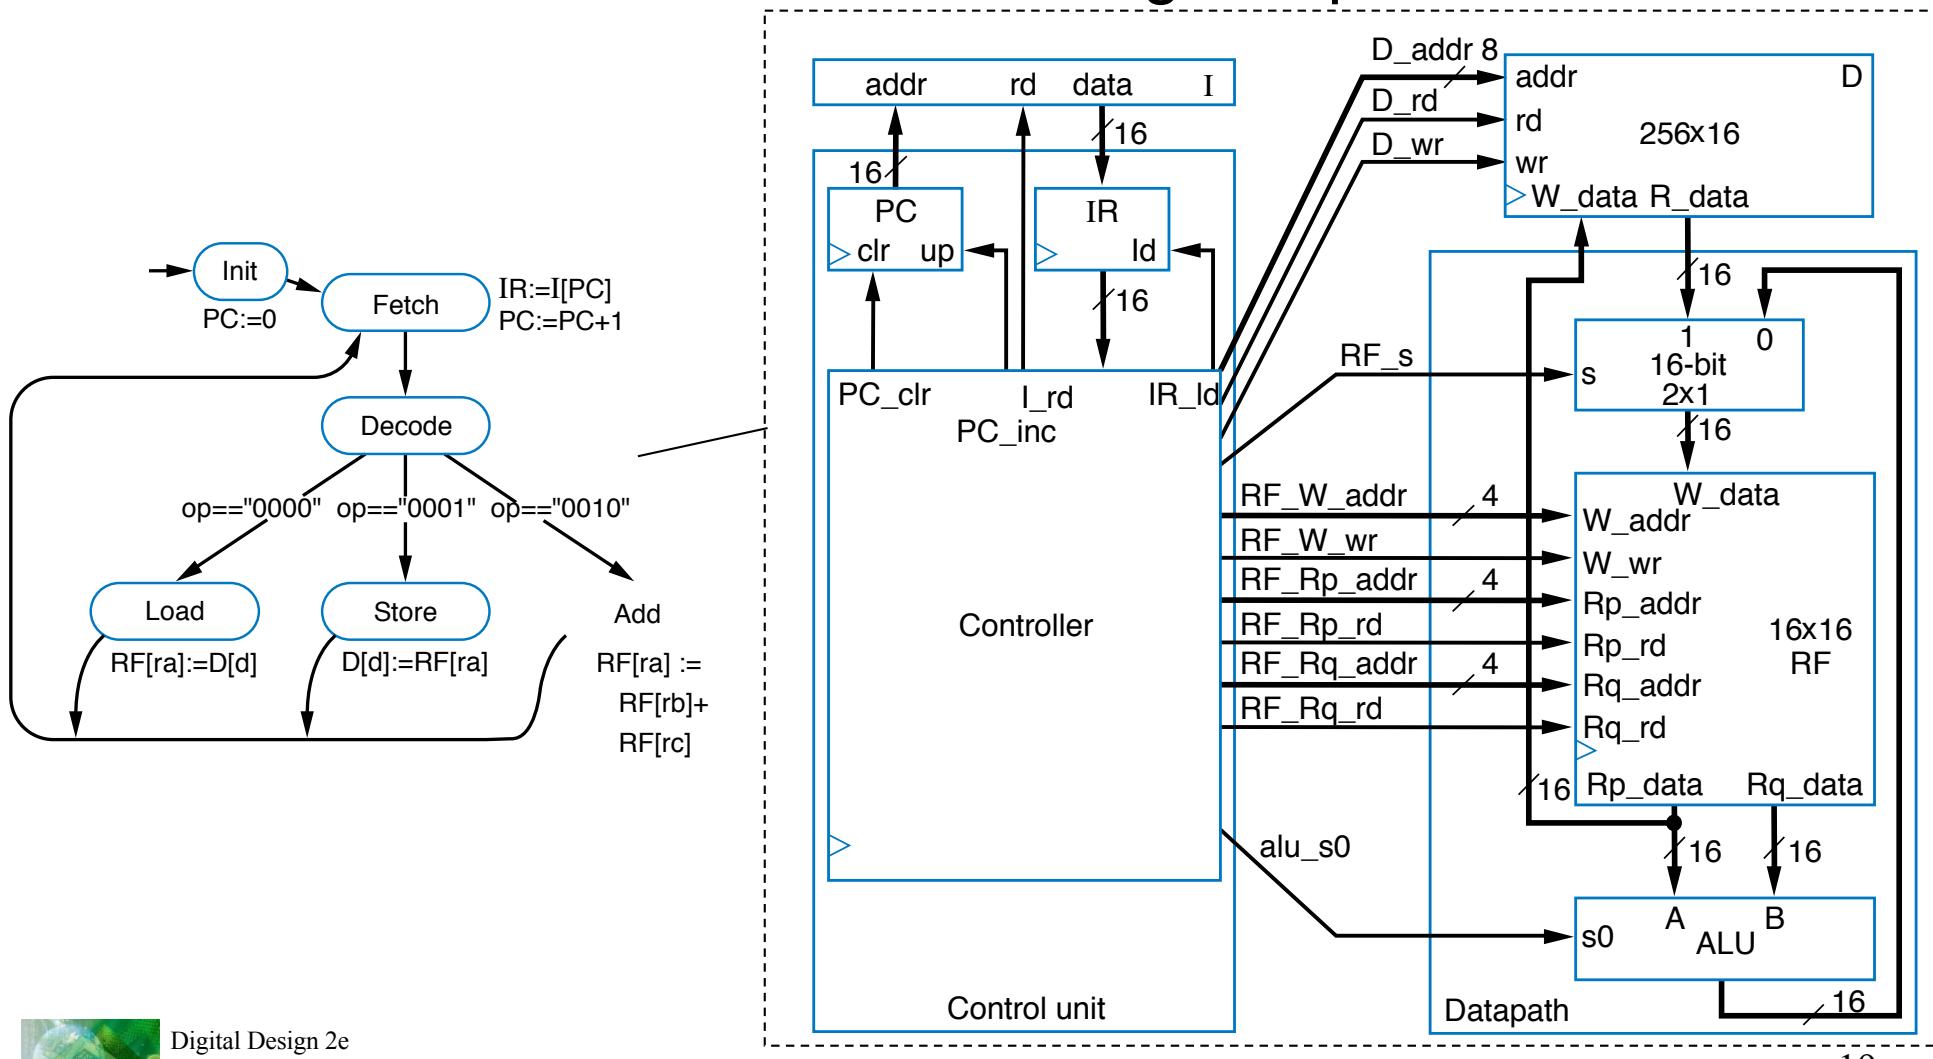
Control-Unit and Datapath for Three-Instruction Processor

- To design the processor, we can begin with a high-level state machine description of the processor's behavior



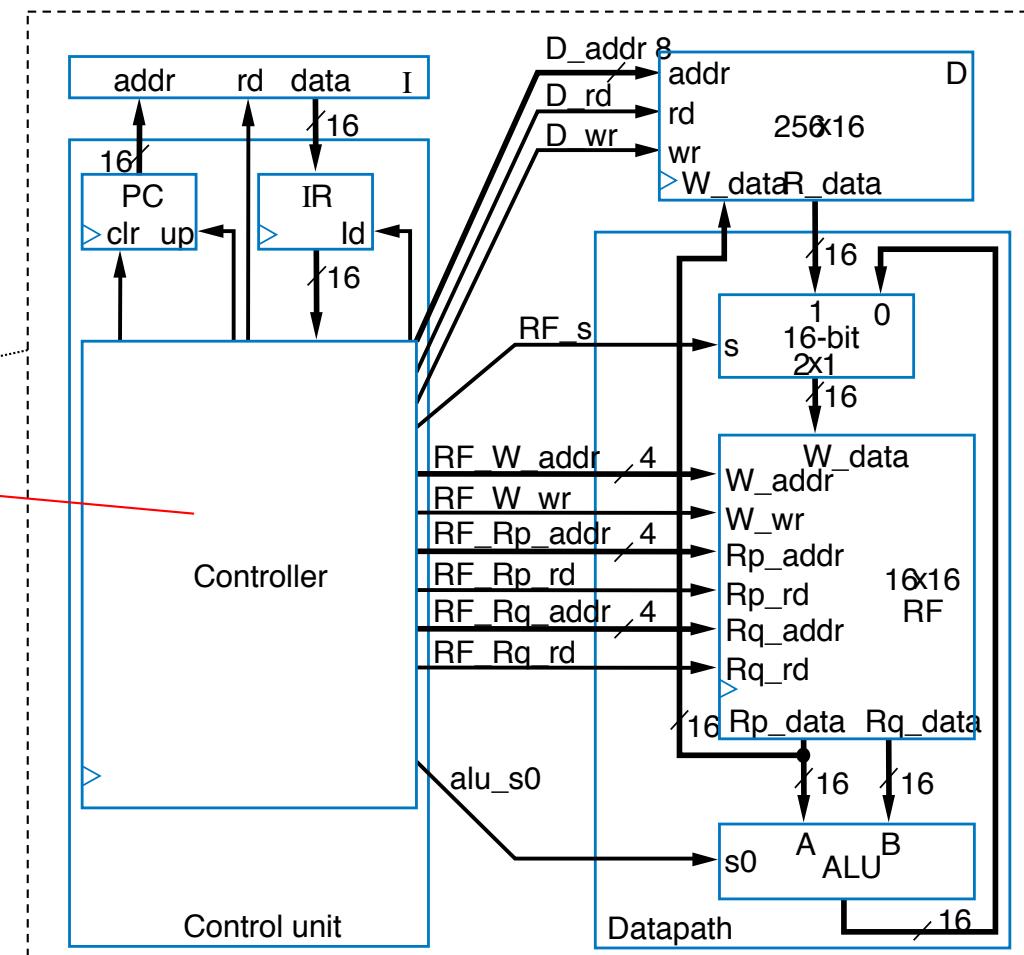
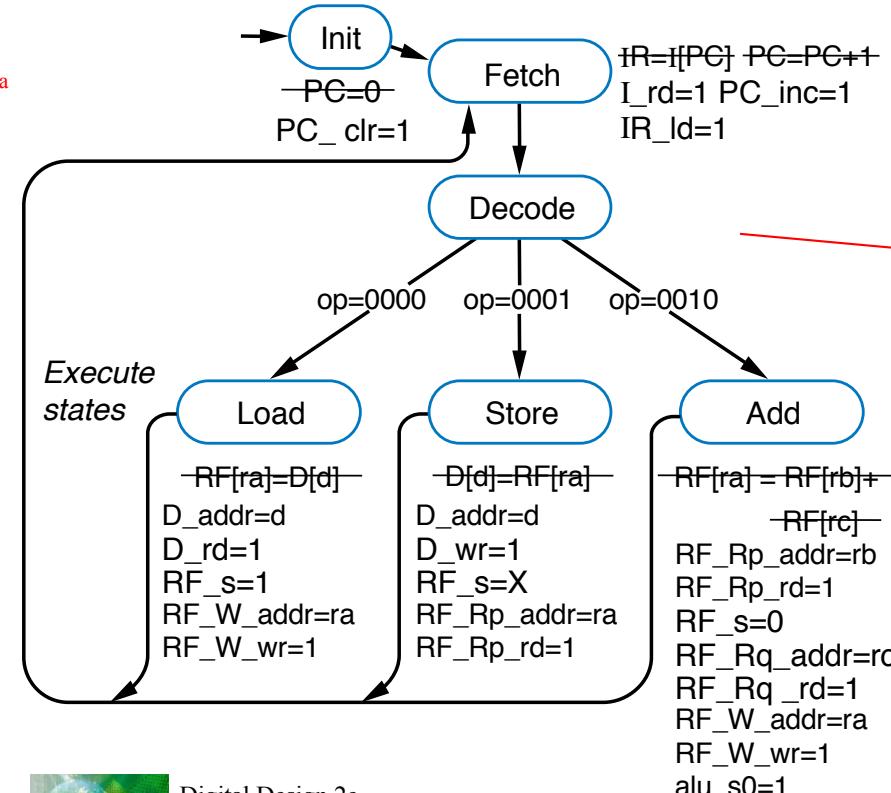
Control-Unit and Datapath for Three-Instruction Processor

- Create detailed connections among components



Control-Unit and Datapath for Three-Instruction Processor

- Convert high-level state machine description of entire processor to FSM description of controller that uses datapath and other components to achieve same behavior



A Six-Instruction Programmable Processor

- Let's add three more instructions:
 - Load-constant* instruction—**0011 r₃r₂r₁r₀ c₇c₆c₅c₄c₃c₂c₁c₀**
 - MOV Ra, #c**—specifies the operation $RF[a] = c$
 - Subtract* instruction—**0100 ra₃ra₂ra₁ra₀ rb₃rb₂rb₁rb₀ rc₃rc₂rc₁rc₀**
 - SUB Ra, Rb, Rc**—specifies the operation $RF[a] = RF[b] - RF[c]$
 - Jump-if-zero* instruction—**0101 ra₃ra₂ra₁ra₀ 0₇0₆0₅0₄0₃0₂0₁0₀**
 - JMPZ Ra, offset**—specifies the operation $PC = PC + \text{offset}$ if $RF[a] = 0$

TABLE 8.1 Six-instruction instruction set..

Instruction	Meaning
MOV Ra, d	$RF[a] = D[d]$
MOV d, Ra	$D[d] = RF[a]$
ADD Ra, Rb, Rc	$RF[a] = RF[b] + RF[c]$
MOV Ra, #C	$RF[a] = C$
SUB Ra, Rb, Rc	$RF[a] = RF[b] - RF[c]$
JMPZ Ra, offset	$PC = PC + \text{offset}$ if $RF[a] = 0$

TABLE 8.2 Instruction opcodes.

Instruction	Opcode
MOV Ra, d	0000
MOV d, Ra	0001
ADD Ra, Rb, Rc	0010
MOV Ra, #C	0011
SUB Ra, Rb, Rc	0100
JMPZ Ra, offset	0101



Extending the Control-Unit and Datapath

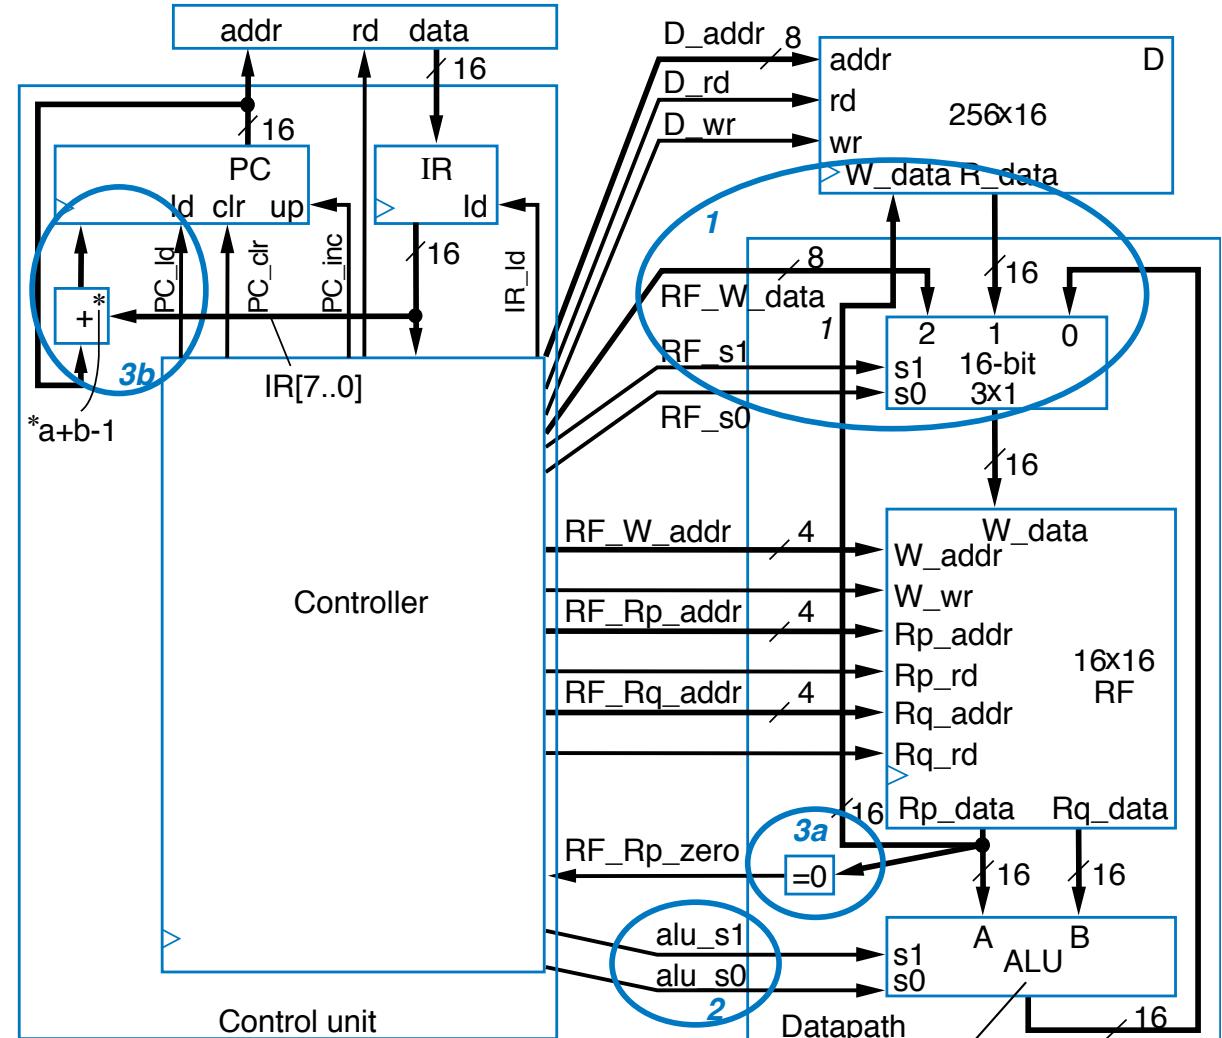
1: The *load constant* instruction requires that the register file be able to load data from $IR[7..0]$, in addition to data from data memory or the ALU output. Thus, we widen the register file's multiplexer from 2×1 to 3×1 , add another mux control signal, and also create a new signal coming from the controller labeled RF_W_data , which will connect with $IR[7..0]$.

2: The subtract instruction requires that we use an ALU capable of subtraction, so we add another ALU control signal.

3: The jump-if-zero instruction requires that we be able to detect if a register is zero, and that we be able to add $IR[7..0]$ to the *PC*.

3a: We insert a datapath component to detect if the register file's *Rp* read port is all zeros (that component would just be a NOR gate).

3b: We also upgrade the *PC* register so it can be loaded with *PC* plus $IR[7..0]$. The adder used for this also subtracts 1 from the sum, to compensate for the fact that the *Fetch* state already added 1 to the *PC*.



s1 s0	ALU operation
0 0	pass A through
0 1	A+B
1 0	A-B



Controller FSM for the Six-Instruction Processor

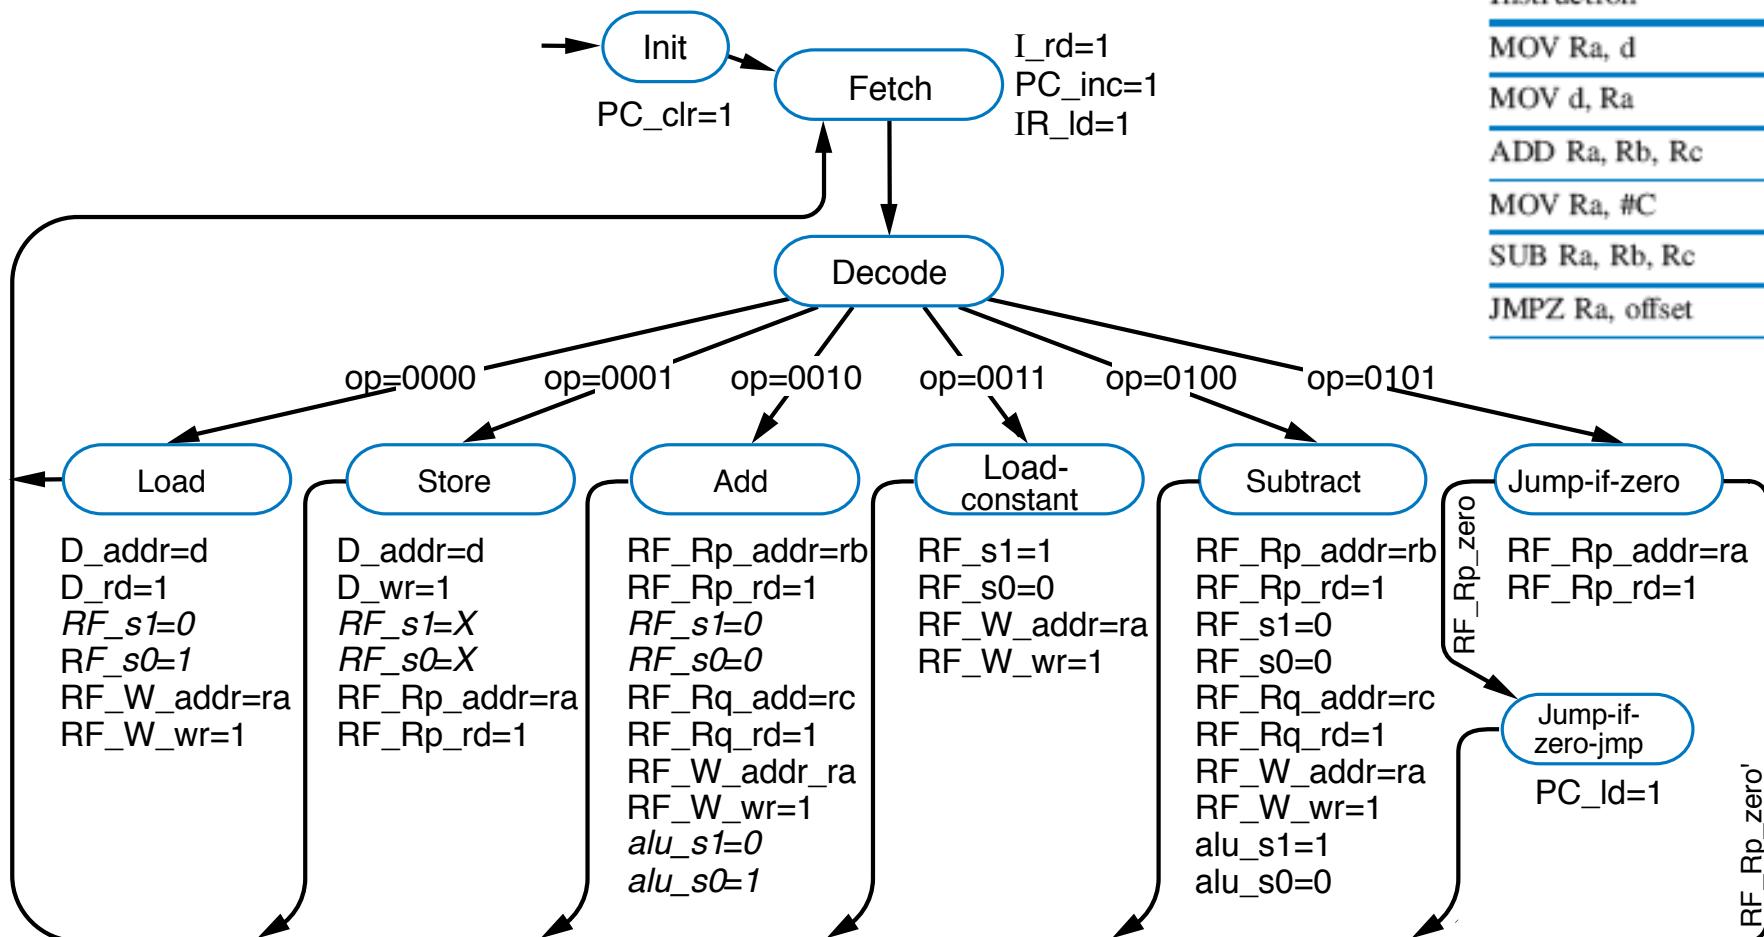


TABLE 8.2 Instruction opcodes.

Instruction	Opcode
MOV Ra, d	0000
MOV d, Ra	0001
ADD Ra, Rb, Rc	0010
MOV Ra, #C	0011
SUB Ra, Rb, Rc	0100
JMPZ Ra, offset	0101



Program for the Six-Instruction Processor

- Example program – Count number of non-zero words in D[4] and D[5]
 - Result will be either 0, 1, or 2
 - Put result in D[9]

MOV R0, #0; // initialize result to 0	0011 0000 00000000
MOV R1, #1; // constant 1 for incrementing result	0011 0001 00000001
MOV R2, 4; // get data memory location 4	0000 0010 00000100
JMPZ R2, lab1; // if zero, skip next instruction	0101 0010 00000010
ADD R0, R0, R1; // not zero, so increment result	0010 0000 0000 0001
lab1:MOV R2, 5; // get data memory location 5	0000 0010 00000101
JMPZ R2, lab2; // if zero, skip next instruction	0101 0010 00000010
ADD R0, R0, R1; //not zero, so increment result	0010 0000 0000 0001
lab2:MOV 9, R0; // store result in data memory location 9	0001 0000 00001001

(a)

TABLE 8.2 Instruction opcodes.

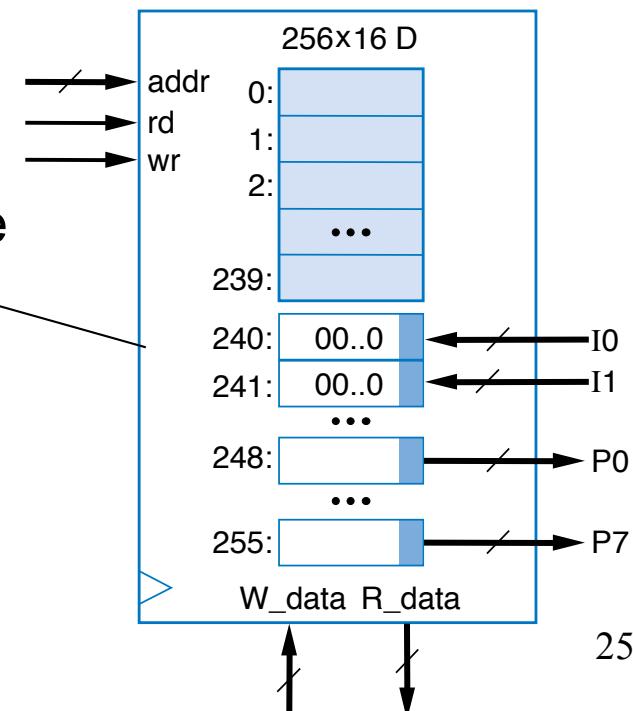
(b)

Instruction	Opcode
MOV Ra, d	0000
MOV d, Ra	0001
ADD Ra, Rb, Rc	0010
MOV Ra, #C	0011
SUB Ra, Rb, Rc	0100
JMPZ Ra, offset	0101



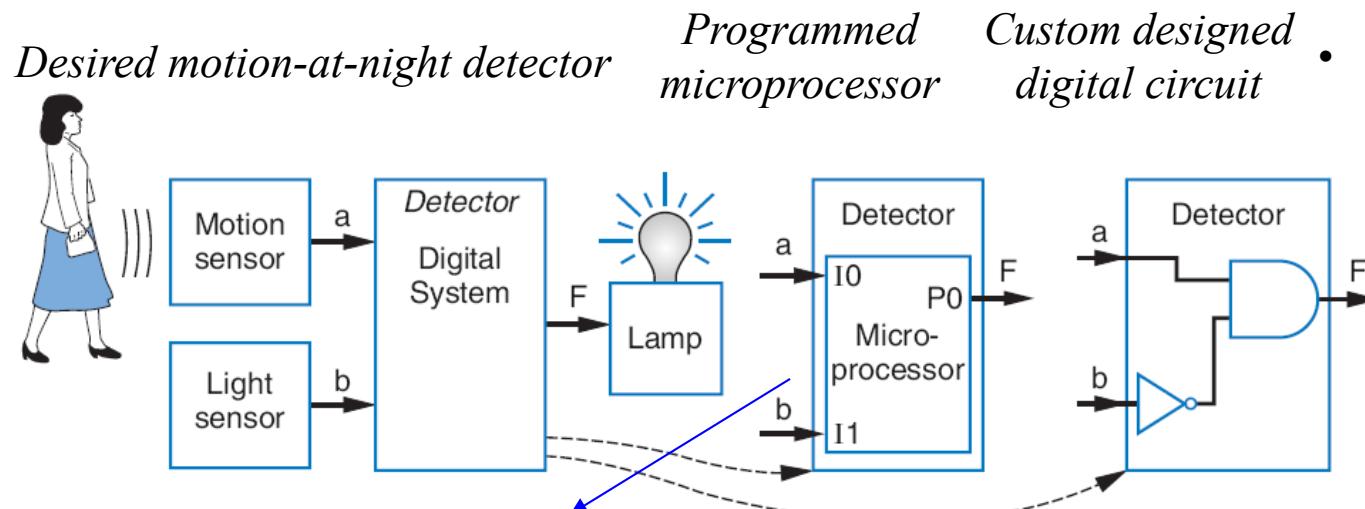
Further Extensions to the Programmable Processor

- Typical processor instruction set will contain dozens of data movement (e.g., loads, stores), ALU (e.g., add, sub), and flow-of-control (e.g., jump) instructions
 - Extending the control-unit/datapath follows similarly to previously-shown extensions
- Input/output extensions
 - Certain memory locations may actually be external pins
 - e.g, D[240] may represent 8-bit input I0, D[255] may represent 8-bit output P7

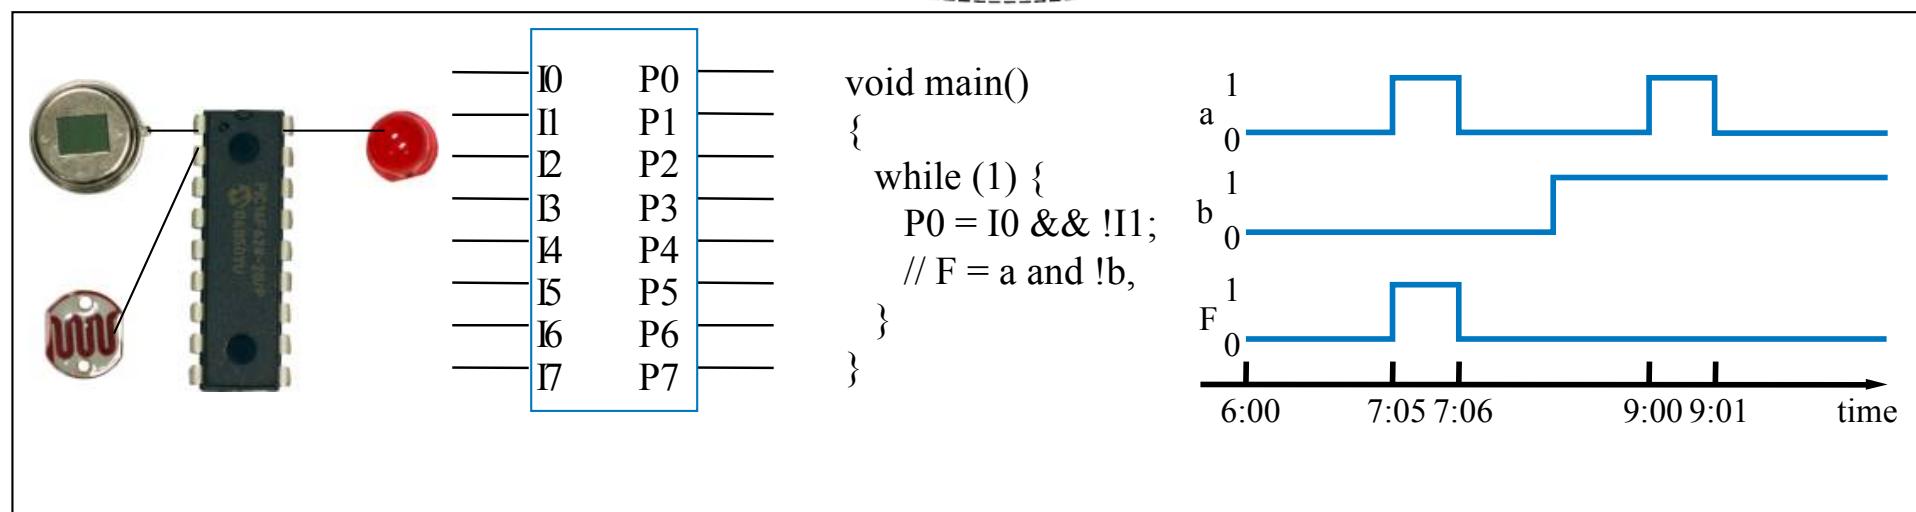


Program using I/O Extensions – Recall Chpt 1

C-Program Example



- Microprocessors a common choice to implement a digital system
 - Easy to program
 - Cheap (as low as \$1)
 - Available now

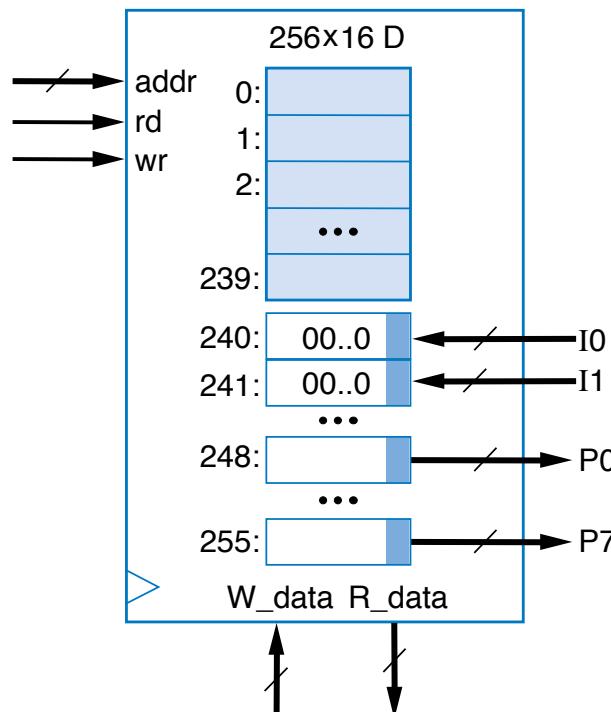


Program Using Input/Output Extensions

Underlying assembly code for C expression $I0 \&& !I1$.

- 0: MOV R0, 240 // move $D[240]$, which is the value at pin $I0$, into $R0$
- 1: MOV R1, 241 // move $D[241]$, which is that value at pin $I1$, into $R1$
- 2: NOT R1, R1 // compute $!I1$, assuming existence of a complement instruction
- 3: AND R0, R0, R1 // compute $I0 \&& !I1$, assuming an AND instruction
- 4: MOV 248, R0 // move result to $D[248]$, which is pin $P0$

```
void main()
{
    while (1) {
        P0 = I0 && !I1;
        // F = a and !b,
    }
}
```



Chapter Summary

- Programmable processors are widely used
 - Easy availability, short design time
- Basic architecture
 - Datapath with register file and ALU
 - Control unit with PC, IR, and controller
 - Memories for instructions and data
 - Control unit fetches, decodes, and executes
- Three-instruction processor with machine-level programs
 - Extended to six instructions
 - Real processors have dozens or hundreds of instructions
 - Extended to access external pins
 - Modern processors are far more sophisticated
- Instructive to see how one general circuit (programmable processor) can execute variety of behaviors just by programming 0s and 1s into an instruction memory

